

QMC integration of non-smooth functions: application to pricing exotic options

Matteo Burzio Luca Civilla Francesco Vinciguerra

January 11, 2026

Abstract

The goal of this project was to estimate the price of Asian call and binary digital Asian options, formulated as high-dimensional integrals. We established baseline convergence rates using Crude Monte Carlo (CMC) and Randomized Quasi-Monte Carlo (RQMC). To restore QMC efficiency for non-smooth integrands, we applied pre-integration smoothing. Finally, we employed importance sampling to achieve variance reduction in the out-of-the-money regime.

Contents

1	Introduction	2
2	Theory and Background	2
2.1	Option pricing framework	2
2.2	Asian options	3
2.3	Numerical Integration Methods	3
2.3.1	Crude Monte Carlo (CMC)	3
2.3.2	Quasi-Monte Carlo (QMC)	3
2.4	Smoothing by pre-integration	4
2.5	Active subspaces and optimal smoothing direction	5
2.6	Sample size in the out of the money regime	6
2.7	Importance Sampling via Optimal Drift (ODIS)	6
2.8	Importance sampling for QMC estimators	7
3	Simulation setup	7
3.1	Model parameters	7
3.2	Performance Metrics and Ground Truth	8
4	Further Analysis and Discussion of Numerical Results	8
4.1	Recovery of QMC Convergence Rates	8
4.2	Synergy between ODIS and Smoothing in the OTM Regime	9
4.3	Efficiency and Computational Trade-offs	9

5 Numerical Results and Graphical Discussion	9
5.1 Analysis of At-The-Money Options ($K = 100$)	9
5.2 Out-of-the-Money Regime and Variance Reduction ($K = 120$)	11
5.3 Sample Size Requirements for Target Accuracy	12
6 Conclusions	12
A Python Code	13
A.1 Arithmetic Asian Option Implementation	13
A.2 Digital Asian Option Implementation	26

1 Introduction

The task of computing high-dimensional integrals is a problem which arises often in the context of financial mathematics, for example when dealing with evaluating Asian options. Due to the absence of closed form solutions, the standard approach is to use numerical methods which don't suffer from the curse of dimensionality, of which Monte Carlo methods are the standard benchmark. Even if they are capable of obtaining a robust convergence rate of $O(N^{-1/2})$, Crude Monte Carlo methods often are too computationally expensive. For this reason, Quasi-Monte Carlo methods are preferred since they theoretically are capable of reaching the superior convergence rate of $O(N^{-1})$ (up to logarithmic factors).

However, the efficacy of QMC is linked to the smoothness of the integrand, which, in the case of the payoff functions for Asian options, presents either a kink or a jump, breaking the regularity requirements for the improved convergence rate. To tackle this, we consider a smoothing technique which involves analytically integrating along a certain direction (in our implementation the specific one chosen is the primary active direction).

Furthermore, we address the problematic high relative variance of the deep out-of-the-money regime by implementing a variance reduction technique known as Optimal Drift Importance Sampling, which shifts the samples closer to the strike price in order to make exercising the option a less rare event.

2 Theory and Background

2.1 Option pricing framework

Let $T > 0$ and S be the price of a stock given by the solution of the stochastic differential equation

$$\begin{cases} dS_t = rS_t dt + \sigma S_t dW_t, & t \in [0, T] \\ S_0 = s_0, \end{cases} \quad (1)$$

where r is the interest rate, σ is the volatility, s_0 is the initial price and $W = (W_t, t \in [0, T])$ is a standard one-dimensional Wiener process. It can be shown that

$$\forall t \in [0, T], \quad S_t = s_0 e^{(r - \frac{\sigma^2}{2})t + \sigma W_t}, \quad (2)$$

which allows us to write the stock price S_t as a function of W_t .

Consider an option with maturity T based on the stock price S .

The payoff $\Psi = \Psi((S_t)_{t \in [0, T]})$ is the function determining the cash flow received by the option holder at maturity T , contingent on the state of the underlying asset S .

The value of an option is defined as the discounted expected payoff under the risk-neutral measure, formally

$$V = e^{-rT} \mathbb{E}[\Psi]. \quad (3)$$

2.2 Asian options

An Asian option is a path-dependent exotic derivative whose payoff is determined by the average price of the underlying asset S , rather than the price at maturity alone (as in European options). In particular, we consider arithmetic Asian options, where the average is taken over a discrete partition of $[0, T]$.

Let $d \in \mathbb{N}$, we define the partition of the time interval $[0, T]$ as $t_i := i\Delta t$ for $i = 1, \dots, d$ where $\Delta t := T/d$. The payoffs of Asian options can then be expressed as a function of the random vector $\mathbf{W} = (W_{t_1}, \dots, W_{t_d})$, which is distributed as a multivariate Gaussian with mean zero and covariance matrix $\Sigma_{ij} = \min(t_i, t_j)$.

Let $K > 0$ be the strike price. We will focus on the two following kinds of Asian options:

- **Asian call option:**

$$\Psi_1(\mathbf{W}) := \left(\frac{1}{d} \sum_{i=1}^d S_{t_i}(W_{t_i}) - K \right)_+ = \left(\frac{1}{d} \sum_{i=1}^d S_{t_i}(W_{t_i}) - K \right) \mathbb{1}_{\{\frac{1}{d} \sum_{i=1}^d S_{t_i}(W_{t_i}) - K \geq 0\}} \quad (4)$$

- **Binary digital Asian option:**

$$\Psi_2(\mathbf{W}) := \mathbb{1}_{\left\{ \frac{1}{d} \sum_{i=1}^d S_{t_i}(W_{t_i}) - K \geq 0 \right\}} \quad (5)$$

The values V_i with $i = 1, 2$ of these Asian options are given by

$$V_i = e^{-rT} \mathbb{E}[\Psi_i(\mathbf{W})] = \frac{e^{-rT}}{(2\pi)^{d/2} \sqrt{\det(C)}} \int_{\mathbb{R}^m} \Psi_i(\mathbf{w}) e^{-\frac{1}{2} \mathbf{w}^T \Sigma^{-1} \mathbf{w}} d\mathbf{w} \quad (6)$$

2.3 Numerical Integration Methods

Since the integral (6) generally lacks a closed-form analytical solution, numerical approximation is necessary. Given the high dimensionality of the problem ($d \gg 1$), deterministic quadrature rules are subject to the curse of dimensionality, making Monte Carlo simulation the preferred methodology.

2.3.1 Crude Monte Carlo (CMC)

The standard Monte Carlo estimator approximates the expected value by the sample mean of N independent and identically distributed (i.i.d.) realizations of \mathbf{W} :

$$\hat{V}_{CMC} = \frac{e^{-rT}}{N} \sum_{i=1}^N \Psi(\mathbf{W}^{(i)}), \quad \mathbf{W}^{(i)} \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}_d(\mathbf{0}, \Sigma) \quad (7)$$

By the Strong Law of Large Numbers (SLLN), \hat{V}_{CMC} converges almost surely to the true value V as $N \rightarrow \infty$. The convergence rate is governed by the Central Limit Theorem (CLT), which states that the error is asymptotically normal with standard deviation proportional to $\sigma N^{-1/2}$, where $\sigma^2 = \text{Var}[\Psi(\mathbf{W})]$. Thus, the convergence rate is $O(N^{-1/2})$, which is robust to dimension but computationally slow.

2.3.2 Quasi-Monte Carlo (QMC)

To accelerate convergence, we consider Quasi-Monte Carlo methods. Instead of random samples, QMC utilizes a deterministic low-discrepancy sequence $\mathcal{S} = \{\mathbf{X}^{(n)}\}_{n \in \mathbb{N}}$ in the unit hypercube $[0, 1]^d$, designed to fill the space more uniformly than random points. The estimator is given by

$$\hat{V}_{QMC} = \frac{e^{-rT}}{N} \sum_{i=1}^N \Psi(A\Phi^{-1}(\mathbf{X}^{(i)})), \quad (8)$$

where Φ^{-1} is the inverse cumulative distribution function of the standard normal applied element-wise, and A is the lower triangular matrix obtained from the Cholesky decomposition of the covariance matrix Σ , satisfying $\Sigma = AA^T$.

The “a priori” error bound for QMC is provided by the Koksma-Hlawka inequality:

$$|V - \hat{V}_{QMC}| \leq \|\Psi\|_{HK} \cdot D^*(\mathcal{S}_N), \quad (9)$$

where $\|\cdot\|_{HK}$ denotes the total variation norm in the sense of Hardy and Krause and D_N^* is the star-discrepancy of the point set. Since low-discrepancy sequences satisfy $D^*(\mathcal{S}_N) = O(N^{-1}(\log N)^d)$, QMC theoretically achieves a convergence rate of $O(N^{-1})$ (up to logarithmic terms), provided the integrand has bounded variation.

2.4 Smoothing by pre-integration

However, the payoff functions for Asian options are non-smooth:

- Ψ_1 has a continuous derivative discontinuity (a “kink”) at the strike K .
- Ψ_2 has a jump discontinuity at K .

These singularities result in unbounded variation in high dimensions, causing the convergence rate of standard QMC to degrade to $O(N^{-1/2})$, providing little advantage over CMC.

In order to maintain the favorable convergence rate of QMC, one can implement the following pre-integration trick.

Consider the integral over \mathbb{R}^d with respect to some distribution $\prod_{j=1}^d \rho_j(x_j)$ of an integrand $f : \mathbb{R}^d \rightarrow \mathbb{R}$ of the form

$$f(\mathbf{x}) = \theta(\mathbf{x}) \mathbb{1}_{\{\phi(\mathbf{x}) \geq 0\}},$$

where θ and ϕ are smooth functions. Also, assume that $\frac{\partial \phi}{\partial x_j}(\mathbf{x}) > 0 \forall \mathbf{x} \in \mathbb{R}^d$ for some $j = 1, \dots, d$ and that $\phi(\mathbf{x}) \rightarrow \pm\infty$ as $x_j \rightarrow \pm\infty$.

The method consists in rewriting the integral in the following way:

$$\int_{\mathbb{R}^d} f(\mathbf{x}) \rho(\mathbf{x}) d\mathbf{x} = \int_{\mathbb{R}^{d-1}} \underbrace{\left(\int_{\mathbb{R}} f(x_j, \mathbf{x}_{-j}) \rho_j(x_j) dx_j \right)}_{:= p(\mathbf{x}_{-j})} \rho_{-j}(\mathbf{x}_{-j}) d\mathbf{x}_{-j}. \quad (10)$$

where $\mathbf{x} = (x_j, \mathbf{x}_{-j})$, with $x_{-j} = (x_1, x_2, \dots, x_{j-1}, x_{j+1}, \dots, x_d)$, and similarly for ρ_j and ρ_{-j} . The reason for doing this is that the resulting inner function $p(\mathbf{x}_{-j})$ is smooth, i.e., it doesn’t have a kink or a jump anymore. Notice that for any fixed $\mathbf{x}_{-j} \in \mathbb{R}^{d-1}$ by the assumptions on the derivatives of ϕ , the function $\mathbf{x}_j \mapsto f(x_j, \mathbf{x}_{-j})$ has a jump at the (unique) point where $\phi(x_j, \mathbf{x}_{-j}) = 0$. It follows from the implicit function theorem that for each \mathbf{x}_{-j} , there exists a unique value $x_j^* := \psi(\mathbf{x}_{-j})$ for which $\phi(x_j, \mathbf{x}_{-j}) < 0$ if $x_j < x_j^*$ and $\phi(x_j, \mathbf{x}_{-j}) > 0$ if $x_j > x_j^*$. Thus, given the chosen form of f , we can write

$$p(\mathbf{x}_{-j}) := \int_{\mathbb{R}} f(x_j, \mathbf{x}_{-j}) \rho_j(x_j) dx_j = \int_{\psi(\mathbf{x}_{-j})}^{+\infty} \theta(x_j, \mathbf{x}_{-j}) \rho_j(x_j) dx_j, \quad (11)$$

with both θ and ψ smooth, and the integral

$$\int_{\mathbb{R}^d} f(\mathbf{x}) \rho(\mathbf{x}) d\mathbf{x} = \int_{\mathbb{R}^{d-1}} p(\mathbf{x}_{-j}) \rho_{-j}(\mathbf{x}_{-j}) d\mathbf{x}_{-j} \quad (12)$$

hopefully now has a smooth integrand.

In the specific case of Asian options, first note that both payoffs (4) and (5) satisfy the regularity conditions required to use this technique.

While the pricing integral V_i in (6) is originally defined under a correlated measure with covariance Σ , we can perform a change of variables using the Cholesky decomposition $\Sigma = AA^T$. Let $\mathbf{Z} \sim \mathcal{N}_d(\mathbf{0}, I_d)$ be a vector of independent standard normal variables such that $\mathbf{W} = A\mathbf{Z}$. The integral (6) can be rewritten as

$$V_i = e^{-rT} \mathbb{E}[\tilde{\Psi}_i(\mathbf{Z})] = e^{-rT} \int_{\mathbb{R}^d} \tilde{\Psi}_i(\mathbf{z}) \rho(\mathbf{z}) d\mathbf{z}, \quad (13)$$

where the new integrand $\tilde{\Psi}_i(\mathbf{z}) := \Psi_i(A\mathbf{z})$ is now defined over a product distribution (since ρ now is the density of a standard multivariate Gaussian), satisfying the requirement for pre-integration.

2.5 Active subspaces and optimal smoothing direction

Standard coordinate-aligned pre-integration may be suboptimal if the discontinuity is not aligned with one of the canonical axes. To maximize the regularizing effect of pre-integration, it is necessary to identify the direction in the input space along which the payoff functions $\tilde{\Psi}_i$ exhibit the greatest variability (in this specific case, a jump or a kink). A possible choice involves building the path via a Brownian Bridge Construction, selecting the first dimension as the smoothing direction because it captures the largest portion of the total path variance. Instead we choose to follow the protocol proposed in [1] which chooses the direction which maximizes the expected squared directional derivative of the payoff, which corresponds to the solution of the following optimization problem

$$\max_{\|\theta\|=1} \mathbb{E}[(\nabla \tilde{\Psi}(\mathbf{Z})^T \theta)^2] = \max_{\|\theta\|=1} \theta^T C \theta, \quad (14)$$

where $\mathbf{Z} \sim \mathcal{N}_d(\mathbf{0}, I_d)$ and $C \in \mathbb{R}^{d \times d}$ defined as $C = \mathbb{E}[\nabla \tilde{\Psi}(\mathbf{Z}) \nabla \tilde{\Psi}(\mathbf{Z})^T]$ is the uncentered gradient covariance matrix. For the binary digital Asian option, where the gradient vanishes almost everywhere, we approximate the active subspace using the gradient of the corresponding Asian call payoff as a smooth proxy.

Since C is symmetric and positive semi-definite, it admits the spectral decomposition $C = U \Lambda U^T$, where $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_d)$ contains the eigenvalues sorted in descending order ($\lambda_1 \geq \dots \geq \lambda_d \geq 0$), and $U = [\mathbf{u}_1, \dots, \mathbf{u}_d]$ contains the corresponding orthonormal eigenvectors.

The leading eigenvector \mathbf{u}_1 defines the primary active direction (i.e., the best direction to perform the pre-integration). In fact, by the Rayleigh-Ritz theorem, the solution of (14) is precisely the leading eigenvector \mathbf{u}_1 , with the maximum value being the corresponding eigenvalue λ_1 . By pre-integrating along \mathbf{u}_1 , we resolve the dimension responsible for the largest fluctuations in the function value, leaving a residual integration problem over the subspace where the function is comparatively flat.

To implement this, we construct an orthogonal matrix $Q \in \mathbb{R}^{d \times d}$ such that its first column aligns with the active direction \mathbf{u}_1 . This is achieved efficiently via a Householder reflection. Let \mathbf{e}_1 denote the first canonical basis vector of \mathbb{R}^d . We define the Householder vector $\mathbf{v} := \mathbf{u}_1 - \|\mathbf{u}_1\| \mathbf{e}_1$ and the resulting matrix as

$$Q = I_d - 2 \frac{\mathbf{v} \mathbf{v}^T}{\|\mathbf{v}\|^2}. \quad (15)$$

By construction, Q is symmetric, orthogonal, and satisfies $Q \mathbf{e}_1 = \mathbf{u}_1$. We then define the rotated random variable $\mathbf{Z}' := Q^T \mathbf{Z}$, which remains distributed as $\mathcal{N}_d(\mathbf{0}, I_d)$. The payoff is re-parameterized again as $\tilde{\Psi}_{rotated}(\mathbf{Z}') := \tilde{\Psi}(\mathbf{Z}')$. Since the first component Z_1 captures the variation along the active direction \mathbf{u}_1 , we apply the pre-integration technique over it:

$$\mathbb{E}[\Psi(\mathbf{W})] = \mathbb{E}[\tilde{\Psi}_{rotated}(\mathbf{Z}')] = \int_{\mathbb{R}^{d-1}} \left(\int_{\mathbb{R}} \tilde{\Psi}(z_1, \mathbf{z}_{-1}) \rho_1(z_1) dz_1 \right) \rho_{-1}(\mathbf{z}_{-1}) d\mathbf{z}_{-1}, \quad (16)$$

where ρ_1 denotes the standard normal density and ρ_{-1} the density of a standard normal Gaussian of dimension $d - 1$.

2.6 Sample size in the out of the money regime

When estimating the value of an option in the “out-of-the-money” (OTM) regime, the probability of the option expiring with a positive payoff is low. Consequently, the Crude Monte Carlo estimator \hat{V}_{CMC} often yields zero for small sample sizes, and its relative error can be prohibitively large.

To determine the sample size N required to achieve a target relative accuracy ε in the mean squared sense, we consider the Root Mean Squared Error (RMSE) criterion:

$$\frac{\sqrt{\mathbb{E}[(\hat{V}_{CMC} - V)^2]}}{V} \leq \varepsilon. \quad (17)$$

Since the CMC estimator is unbiased, the Mean Squared Error (MSE) is equal to the variance of the estimator, $\text{Var}[\hat{V}_{CMC}] = \frac{1}{N} \text{Var}[\Psi(\mathbf{W})]$. The condition becomes:

$$\frac{\sqrt{\text{Var}[\Psi(\mathbf{W})]}}{\sqrt{N} \cdot V} \leq \varepsilon. \quad (18)$$

Solving for N , we obtain the required sample size in terms of the coefficient of variation of the payoff:

$$N \geq \left\lceil \frac{\text{Var}[\Psi(\mathbf{W})]}{\varepsilon^2 V^2} \right\rceil. \quad (19)$$

Since the true variance $\text{Var}[\Psi(\mathbf{W})]$ and the true option value V are unknown a priori, they must be estimated numerically.

Equation (19) highlights a critical dependency: for a fixed relative accuracy ε , the required computational effort scales with the variance of the payoff. In the deep OTM regime, the coefficient of variation is typically very large, rendering the CMC approach computationally prohibitive as N grows uncontrollably. To achieve the target precision within a reasonable computational budget, increasing the sample size is often not a viable strategy. Instead, we must focus on reducing the numerator in (19). This motivates the implementation of variance reduction techniques.

2.7 Importance Sampling via Optimal Drift (ODIS)

Importance Sampling (IS) is a variance reduction technique that is particularly effective for estimating probabilities of rare events, which is precisely the case for the payoff of deep out-of-the-money (OTM) options. IS introduces a proposal density $q(\mathbf{w})$ to rewrite the value of the option (6) as

$$V_i = e^{-rT} \int_{\mathbb{R}^d} \tilde{\Psi}_i(\mathbf{z}) \frac{\rho(\mathbf{z})}{q(\mathbf{z})} q(\mathbf{z}) d\mathbf{z} = e^{-rT} \int_{\mathbb{R}^d} \tilde{\Psi}_i(\mathbf{z}) w(\mathbf{z}) q(\mathbf{z}) d\mathbf{z} = e^{-rT} \mathbb{E}_q [\tilde{\Psi}(\mathbf{Z}) w(\mathbf{Z})], \quad (20)$$

where $\mathbf{Z} \sim \mathcal{N}_d(\mathbf{0}, I_d)$, $w = \frac{\rho}{q}$ and the expectation \mathbb{E}_q is taken with respect to the density q .

While the theoretical optimal estimator is achieved by choosing $q^*(\mathbf{z}) \propto |\tilde{\Psi}_i(\mathbf{z})| \rho(\mathbf{z})$, this density is generally intractable as its normalization constant is precisely the integral V_i we wish to compute. Consequently, we restrict our search to a parametric family of distributions that approximates q^* while remaining easy to sample from.

Optimal Drift Importance Sampling (ODIS) employs a mean-shift strategy, restricting q to the family of multivariate Gaussian densities with a shifted mean $\boldsymbol{\mu}$ and identity covariance:

$$q(\mathbf{z}; \boldsymbol{\mu}) = (2\pi)^{-d/2} e^{-\frac{1}{2}\|\mathbf{z}-\boldsymbol{\mu}\|^2}. \quad (21)$$

Under this change of measure, the likelihood ratio simplifies to the exponential of a linear function of \mathbf{z} :

$$w(\mathbf{z}; \boldsymbol{\mu}) = \frac{e^{-\frac{1}{2}\|\mathbf{z}\|^2}}{e^{-\frac{1}{2}\|\mathbf{z}-\boldsymbol{\mu}\|^2}} = e^{-\boldsymbol{\mu}^T \mathbf{z} + \frac{1}{2}\|\boldsymbol{\mu}\|^2}. \quad (22)$$

To select the optimal parameter μ^* , we aim to match the mode of the proposal density $q(\mathbf{z}; \mu)$ to the mode of the optimal zero-variance density $q^*(\mathbf{z})$. This is equivalent to finding the point \mathbf{z} that maximizes the integrand product $\tilde{\Psi}_i(\mathbf{z})\rho(\mathbf{z})$, which corresponds to solving the following optimization problem:

$$\mu^* = \arg \max_{\mathbf{z} \in \mathbb{R}^d} \left(\ln \tilde{\Psi}_i(\mathbf{z}) - \frac{1}{2} \|\mathbf{z}\|^2 \right), \quad (23)$$

where we have taken the logarithm of the product (which doesn't affect the solution since it's an increasing transformation) and ignored the terms not depending on \mathbf{z} . In the case of the binary digital Asian option, where the payoff is an indicator function, the optimization objective is modified to maximize the proposal density at the boundary surface defined by the strike price.

2.8 Importance sampling for QMC estimators

The integration of Importance Sampling within a Quasi-Monte Carlo framework requires mapping the deterministic low-discrepancy sequence to the proposal distribution q rather than the nominal distribution ρ . Let $\mathcal{S} = \{\mathbf{X}^{(n)}\}_{n \in \mathbb{N}}$ be a low-discrepancy sequence in the unit hypercube $[0, 1]^d$. The RQMC-IS estimator is given by:

$$\hat{V}_{RQMC}^{IS} = \frac{e^{-rT}}{N} \sum_{i=1}^N \tilde{\Psi}(G^{-1}(\mathbf{X}^{(i)})) \frac{\rho(G^{-1}(\mathbf{X}^{(i)}))}{q(G^{-1}(\mathbf{X}^{(i)}))}, \quad (24)$$

where $G^{-1} : [0, 1]^d \rightarrow \mathbb{R}^d$ denotes the inverse CDF applied component-wise associated with the proposal density q .

In the context of ODIS, the proposal density $q(\mathbf{z}; \mu)$ corresponds to a mean-shifted multivariate Gaussian $\mathcal{N}_d(\mu, I_d)$. The generation of sample points from this distribution using QMC is computationally direct; we simply shift the standard normal inverse mapping:

$$\mathbf{Z}^{(i)} = \Phi^{-1}(A\mathbf{X}^{(i)}) + \mu, \quad (25)$$

where as before Φ^{-1} is the inverse cumulative distribution function of the standard normal applied element-wise, and A is the lower triangular matrix obtained from the Cholesky decomposition of the covariance matrix Σ , satisfying $\Sigma = AA^T$.

Substituting this transformation and the likelihood ratio $w(\mathbf{z}; \mu) = e^{-\mu^T \mathbf{z} + \frac{1}{2} \|\mu\|^2}$ into the general estimator, we obtain the ODIS-RQMC estimator:

$$\hat{V}_{RQMC}^{ODIS} = \frac{e^{-rT}}{N} \sum_{i=1}^N \tilde{\Psi}(\mathbf{Z}^{(i)}) \exp \left(-\mu^T \mathbf{Z}^{(i)} + \frac{1}{2} \|\mu\|^2 \right). \quad (26)$$

While the effectiveness of CMC with IS is governed by the variance (ℓ_2 norm) of the weighted integrand, the convergence rate of QMC is determined by the total variation in the sense of Hardy and Krause ($\|\cdot\|_{HK}$). A potential concern is that the change of measure might introduce boundary singularities or unbounded variation. However, recent theoretical results [2] demonstrate that for Gaussian proposals the weighted integrand satisfies the necessary light-tailed conditions. Consequently, the combination of RQMC and ODIS is theoretically justified.

3 Simulation setup

3.1 Model parameters

The model parameters are set as follows:

- Initial asset price: $S_0 = 100$

- Risk-free interest rate: $r = 0.1$
- Volatility: $\sigma = 0.1$
- Maturity: $T = 1$
- Discretization: $d = 32, 64, 128, 256, 512$

We evaluate the performance of our estimators for two distinct strike prices:

1. $K = 100$ at-the-money (ATM) scenario ($K = S_0$).
2. $K = 120$ deep out-of-the-money (OTM) scenario ($K > S_0$)

3.2 Performance Metrics and Ground Truth

To assess the accuracy and convergence of the numerical methods, we compute the Root Mean Squared Error (RMSE). Since arithmetic Asian options lack a closed-form analytical pricing formula, we generate a high-precision benchmark value V_{ref} to serve as the ground truth. This reference value is computed using our most robust estimator (RQMC with pre-integration) using a large sample size of $N = 2^{17}$ (131,072 points). For the out-of-the-money case ($K = 120$), this benchmark is further stabilized using the optimal drift (ODIS) as well.

For a given estimator \hat{V} and sample size N , the RMSE is estimated empirically over $R = 50$ independent simulation runs:

$$\text{RMSE}(N) = \sqrt{\frac{1}{R} \sum_{r=1}^R (\hat{V}_N^{(r)} - V_{\text{ref}})^2}, \quad (27)$$

where $\hat{V}_N^{(r)}$ denotes the price estimate obtained in the r -th repetition. This experimental setup, which relies on a high-precision numerical benchmark and independent repetitions to estimate convergence, follows the methodology adopted in [1, 2].

4 Further Analysis and Discussion of Numerical Results

The numerical experiments clearly validate the theoretical framework developed in the preceding sections. Below, we provide a detailed commentary on the performance of the various estimators.

4.1 Recovery of QMC Convergence Rates

The theoretical primary advantage of Randomized Quasi-Monte Carlo (RQMC) is its potential to achieve a convergence rate near $O(N^{-1.0})$. However, as observed in our baseline simulations (Plain RQMC), the lack of smoothness of payoffs, like the "kink" in arithmetic options and the "jump" in digital options, initially pushes this rate toward the $O(N^{-0.5})$ characteristic of Crude Monte Carlo (CMC).

- **Arithmetic Options:** The Plain RQMC achieves a sub-optimal rate (approx. $N^{-0.7}$ to $N^{-0.8}$) due to the derivative discontinuity. The implementation of **Pre-integration along the Active Subspace (AS)** successfully "smooths" this kink. As shown in the convergence plots for all dimensions ($d = 32$ to $d = 512$), the Pre-Int (AS) estimator restores the superior $O(N^{-1.0})$ rate for $d = 32$, significantly outperforming other estimators.
- **Digital Options:** The challenge is more acute here due to the jump discontinuity. Our results confirm that Plain RQMC offers almost no benefit over CMC in this case. The application of the **Pre-integration trick** is crucial: by analytically integrating the jump into a smooth Gaussian tail probability, the RMSE drops by several orders of magnitude, and the convergence rate is restored to near-ideal QMC levels.

4.2 Synergy between ODIS and Smoothing in the OTM Regime

In the out-of-the-money (OTM) regime ($K = 120$), the scarcity of "in-the-money" samples renders standard estimators highly inefficient. The results for $K = 120$ demonstrate a significant joint impact of Importance Sampling and Pre-integration:

- **Impact of ODIS:** The **Optimal Drift Importance Sampling (ODIS)** shift ensures that the simulation focuses on the region near the exercise boundary. In our "Impact of Variance Reduction" plots, the shift from dashed lines (base methods) to solid lines (ODIS-enhanced) represents the variance reduction factor. Our results show a reduction in the RMSE with no significant improvement in the convergence rate, matching the results obtained in [2].
- **The Optimal Estimator:** The combination **Pre-Int (AS) + ODIS** emerges as the most robust solution. While ODIS reduces the variance by shifting the measure, Pre-integration ensures the weighted integrand remains smooth for the RQMC points. This is particularly evident in high dimensions ($d = 512$), where this combination obtains better convergence rate and lower RMSE whereas other methods suffer from the curse of dimensionality.

4.3 Efficiency and Computational Trade-offs

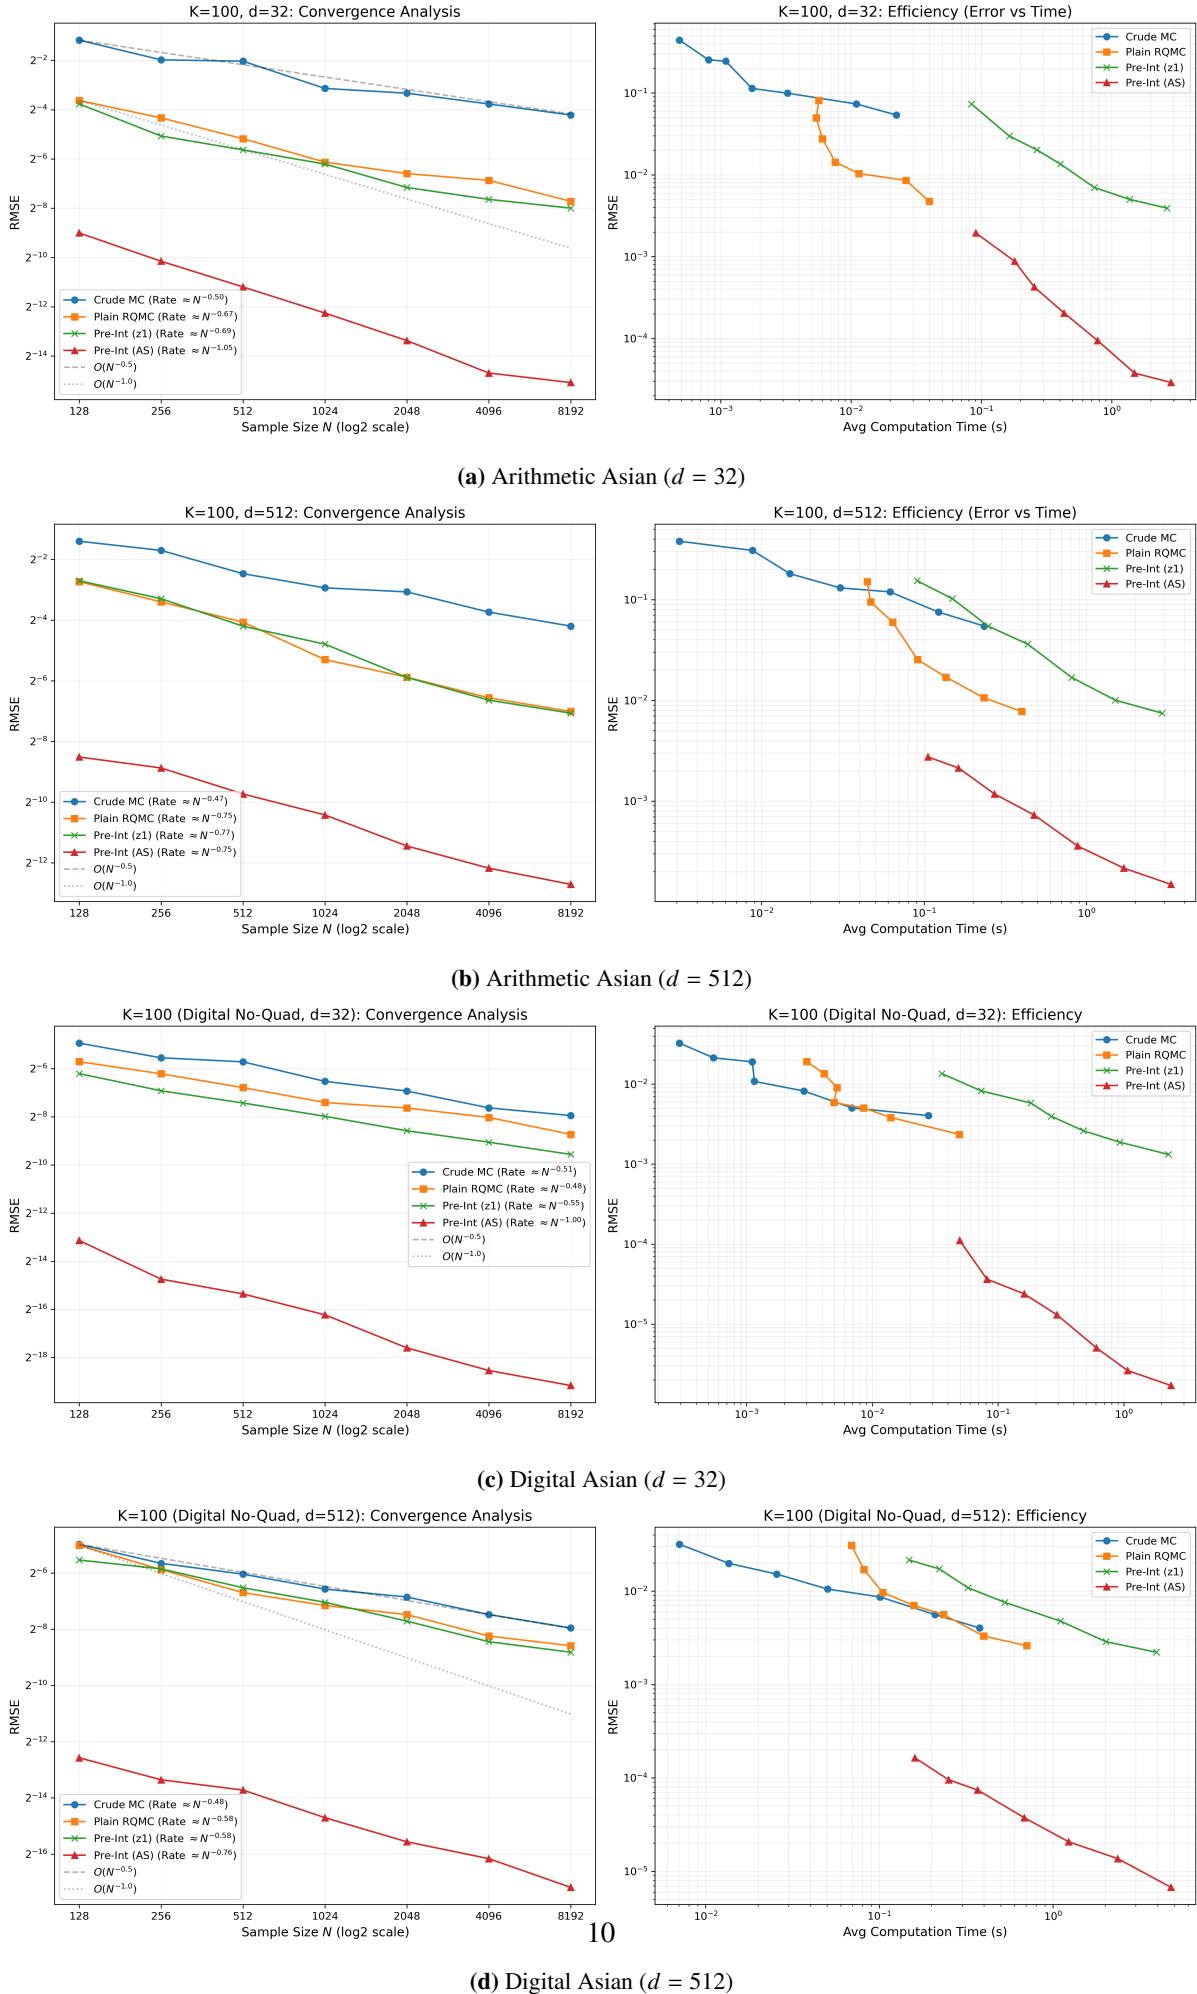
The "Efficiency (Error vs Time)" plots provide a pragmatic perspective. Although the advanced estimators (Pre-Int and ODIS) are slow mainly due to the root-finding (v^*) and gradient covariance matrix (C) calculations, the faster convergence rate makes them a better choice when needing a low tolerance. To achieve a target RMSE of 10^{-2} in the digital case, Crude MC would require millions of samples, whereas **Pre-Int (AS) + ODIS** reaches this threshold in a fraction of a second, proving to be the only viable approach for high-precision exotic option pricing.

5 Numerical Results and Graphical Discussion

In this section, we present the numerical results obtained from our simulations, comparing the performance of the estimators across different dimensions (d) and strike prices (K).

5.1 Analysis of At-The-Money Options ($K = 100$)

The following figures illustrate the convergence and efficiency for the arithmetic and digital cases. For the digital option, the presence of a jump discontinuity at the strike price makes standard integration particularly challenging.



Commentary: As predicted by the theory, the **Crude Monte Carlo** (blue line) consistently follows the $O(N^{-0.5})$ rate. For the **Arithmetic Option**, the "kink" in the payoff degrades the Plain RQMC performance to approximately $N^{-0.67}$. However, the **Pre-Int (AS)** method (red triangles) successfully restores the $O(N^{-1.0})$ convergence in the $d=32$ case. Indeed, Chen et al. [2] implement this technique with at most 80 dimensions. Our work extends this validation to significantly higher dimensions; in such cases the implemented techniques are not able to recover the $O(N^{-1.0})$ rate. We believe that this might be due to the fact that the variance cannot be well explained by looking (in this case integrating) at only one dimension. In the **Digital Case**, the impact of smoothing is even more evident. Without pre-integration, the jump discontinuity causes RQMC to converge as slowly as CMC. By applying the Gaussian smoothing trick (Pre-Int AS), we achieve a convergence rate of N^{-1} , effectively neutralizing the degrading impact of the binary payoff.

5.2 Out-of-the-Money Regime and Variance Reduction ($K = 120$)

For $K = 120$, the probability of exercising the option is low, making the "out-of-the-money" regime a rare-event simulation problem.

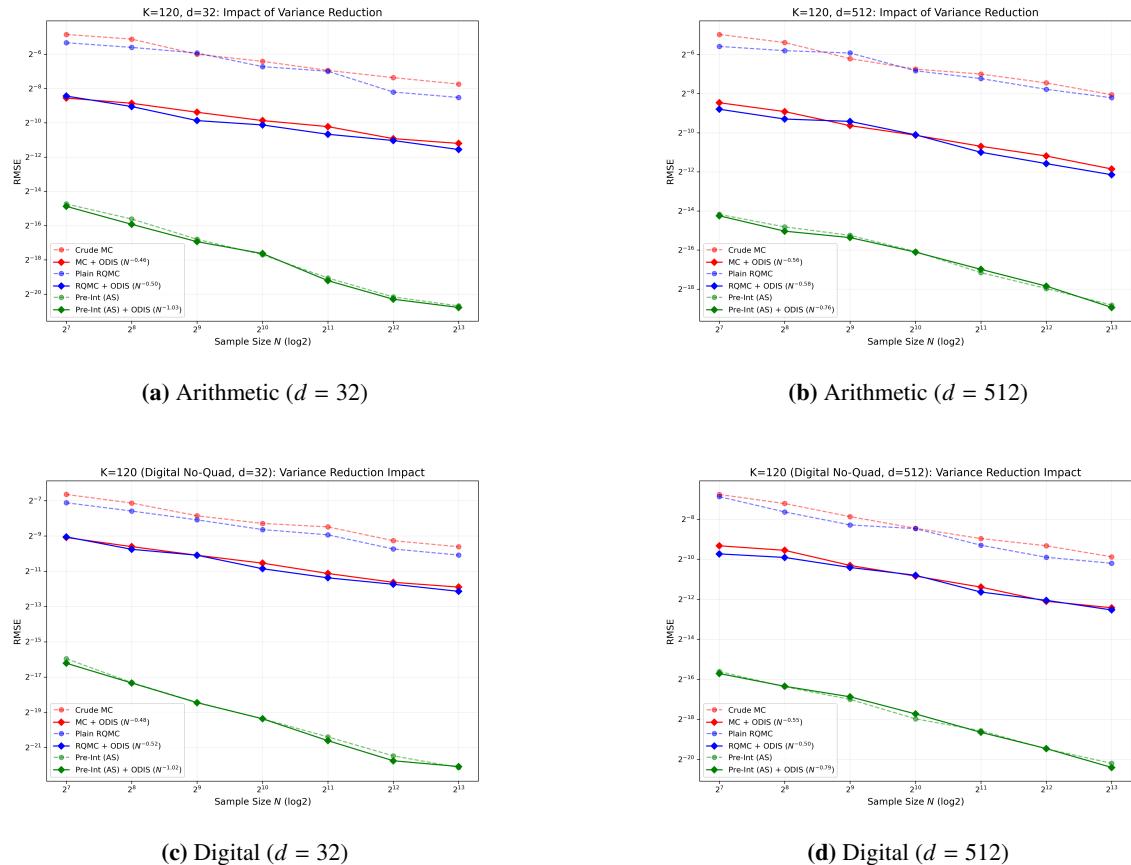


Figure 2: Impact of Variance Reduction for $K = 120$. Comparison across dimensions $d = 32$ and $d = 512$ for both Arithmetic and Digital options. The synergy between ODIS shift and AS Pre-integration consistently leads to the lowest RMSE values (green lines).

Commentary: The plots highlight the critical role of **Optimal Drift Importance Sampling (ODIS)**.

The most efficient estimator is the **Pre-Int (AS) + ODIS**. It combines the dimensionality reduction of the active subspace with the rare-event optimization of ODIS.

5.3 Sample Size Requirements for Target Accuracy

To conclude our numerical analysis, we report the sample size N required to achieve a relative error tolerance of $\varepsilon = 0.01$ in the OTM regime ($K = 120$) using CMC. This metric serves as a proxy for the variance of the estimators and the inherent difficulty of the pricing problem.

Table 1: Required sample size N for 1% relative accuracy ($K = 120$).

Dimension (d)	Digital Option (N)	Arithmetic Option (N)
32	2,640,164	4,996,001
64	2,927,043	5,501,541
128	3,029,561	5,798,022
256	3,103,982	5,901,949
512	3,164,820	5,940,186

The results in Table 1 show that the Arithmetic option consistently requires more samples than the Digital one, likely due to the variability of the continuous payoff compared to the binary indicator. Most importantly, it highlights that while a standard Monte Carlo would require up to 6 million paths, our proposed Pre-Int (AS) + ODIS method achieves superior precision with a fraction of the computational budget.

6 Conclusions

In this project, we investigated the challenges of pricing high-dimensional Asian options characterized by non-smooth payoff functions. Our results confirm that advanced regularization and variance reduction techniques are indispensable for overcoming the limitations of standard Monte Carlo methods.

In summary, the implementation of **pre-integration along the Active Subspace (AS)** successfully addressed the difficulties of "kinks" and "jumps" in the payoffs. This approach effectively restored the ideal QMC convergence rate of $O(N^{-1})$ in the lower dimensional setting, while it improved significantly the convergence rate in the other cases. Furthermore, in the out-of-the-money regime ($K = 120$), the integration of **Optimal Drift Importance Sampling (ODIS)** provided a variance reduction factor of an average order of 10^6 .

As highlighted by our adaptive search, a traditional Crude Monte Carlo approach would require up to 6 million simulations to reach a 1% relative accuracy for $K = 120$. In contrast, the proposed **Pre-Int (AS) + ODIS** framework achieves superior precision with a drastically lower computational budget. Ultimately, this project demonstrates that the synergy between analytical smoothing and optimized sampling represents a robust and scalable solution for modern computational finance.

To improve stability in higher dimension we propose to consider for the Pre-Integration trick a multi-dimensional active subspace.

References

- [1] Sifan Liu and Art B Owen. Pre-integration via active subspaces. *arXiv preprint arXiv:2202.02682*, 2022.
- [2] Jianlong Chen, Yu Xu, Jiarui Du, and Xiaoqun Wang. Randomized quasi-monte carlo and importance sampling for super-fast growing functions with applications to finance. *arXiv preprint arXiv:2510.06705*, 2025.

A Python Code

A.1 Arithmetic Asian Option Implementation

```

1 import os
2 import time
3 import numpy as np
4 import pandas as pd
5 import scipy.stats as stats
6 from scipy.stats import qmc
7 from scipy.optimize import minimize, brentq
8 import matplotlib.pyplot as plt
9 import matplotlib.ticker as mticker
10 from tqdm import tqdm
11
12 np.random.seed(42)
13
14 """
15 This module implements Monte Carlo and Quasi-Monte Carlo simulations for pricing arithmetic Asian
16 options.
17
18 Arithmetic Asian options have payoffs based on the arithmetic average of the underlying asset prices
19 over time.
20
21 The code uses various variance reduction techniques including:
22 - Quasi-Monte Carlo (RQMC) with Sobol sequences
23 - Importance sampling (ODIS - Optimal Drift for Importance Sampling)
24 - Pre-integration estimators using closed-form solutions
25 - Active subspace methods for dimension reduction
26
27 The simulations are performed for different dimensions (number of time steps) and strike prices,
28 with convergence analysis and efficiency comparisons.
29
30 """
31 # =====
32 # Setup and Configuration
33 # =====
34 # Model Parameters
35 S0 = 100 # Initial stock price
36 T = 1.0 # Time to maturity (in years)
37 r = 0.1 # Risk-free interest rate
38 sigma = 0.1 # Volatility of the stock
39
40 # Dimensions to test (number of time steps in the Asian option)
41 DIMENSIONS = [32, 64, 128, 256, 512]
42
43 # Simulation Parameters
44 POWERS = np.arange(7, 14) # Powers from 7 to 13, so sample sizes 2^7 to 2^13
45 SAMPLE_SIZES = 2**POWERS # Actual sample sizes: 128, 256, ..., 8192
46 N_REPEATS = 50 # Number of repetitions for each simulation to compute stable RMSE
47
48 # =====
49 # Adaptive N Finder
50 # =====
51
52 def find_N_specific_arithmetic(model, d, tol=0.01):

```

```

53     """
54     find_n_adaptive(model, ..., tol): Implements the stopping criterion derived in Eq. (19). Rather
55     than
56     fixing N a priori, this routine increases the sample size in batches, monitoring the standard
57     error of
58     the estimator. It terminates the simulation only when the relative RMSE falls below the specified
59     tolerance (e.g., =  $10^{-2}$ ), ensuring the reliability of the OTM price estimates.
60
61     Parameters:
62     - model: An instance of AsianOptionCholesky representing the option model.
63     - d: Dimension (number of time steps).
64     - tol: Tolerance for relative error (default 0.01, i.e., 1%).
65
66     Returns:
67     - N: The required sample size.
68     - mu: The estimated mean payoff.
69     """
70
71     z_score = 1.96 # 95% Confidence interval z-score
72
73     N = 128 # Starting sample size
74     # Generate initial batch of standard normal random variables
75     z_init = np.random.standard_normal((N, model.d))
76     Z_values = model.payoff(z_init) # Compute payoffs for initial batch
77
78     # Initial statistics: mean and variance
79     mu = np.mean(Z_values)
80     var = np.var(Z_values, ddof=1) # Sample variance with ddof=1
81
82     print(f"\n--- Specific Recursive Search (Arithmetic, d={d}, Start N={N}) ---")
83     print(f"{'N':<10} | {'Mean':<10} | {'Var':<10} | {'RelErr':<10}")
84     print("-" * 55)
85
86     while True:
87         # Generate one new standard normal sample
88         z_new = np.random.standard_normal((1, model.d))
89         Z_next = model.payoff(z_new)[0] # Payoff for the new sample
90
91         # Save previous mean for variance update
92         mu_old = mu
93
94         # Recursive update of mean and variance
95         mu = (N / (N + 1)) * mu + Z_next / (N + 1)
96         var = ((N - 1) / N) * var + (Z_next - mu_old)**2 / (N + 1)
97
98         N += 1 # Increment sample size
99
100        # Check for convergence
101        sigma_est = np.sqrt(max(var, 0.0)) # Estimated standard deviation
102
103        if abs(mu) > 1e-12 and sigma_est > 0:
104            half_width = z_score * sigma_est / np.sqrt(N) # Half-width of confidence interval
105            rel_err = half_width / abs(mu) # Relative error
106
107            if rel_err <= tol:
108                print("-" * 55)
109                print(f"Converged at N = {N} for d={d}")
110                print(f"Final Mean: {mu:.6f}")
111                print(f"Final Var: {var:.6f}")

```

```

109         return N, mu
110
111     # Print progress every 5000 samples
112     if N % 5000 == 0:
113         print(f"{N:<10} | {mu:<10.4f} | {var:<10.4f} | {rel_err:<10.2%}")
114
115 # =====
116 # Model: Asian Option with Cholesky Path
117 # =====
118
119 class AsianOptionCholesky:
120     """
121     Represents an arithmetic Asian option using Cholesky decomposition for path simulation.
122
123     This class models the underlying asset paths using correlated Brownian motions,
124     discretized into 'd' time steps. The payoff is based on the arithmetic average
125     of the stock prices at these time steps.
126
127     Attributes:
128     - S0: Initial stock price
129     - K: Strike price
130     - T: Time to maturity
131     - r: Risk-free rate
132     - sigma: Volatility
133     - d: Number of time steps
134     - dt: Time step size
135     - A: Cholesky matrix for correlation
136     - drift_term: Precomputed drift term for each time step
137     """
138     def __init__(self, S0=100, K=100, T=1.0, r=0.1, sigma=0.1, d=32):
139         """
140         create_model(S0, K, T, r, , m): This initialization routine constructs the discrete time grid
141             t_i = i*T/m where i = 1, ..., m and the associated covariance matrix R^{mm} with entries
142             _{ij} = min(t_i, t_j). Crucially, it performs the Cholesky decomposition = A A^T (via
143             numpy.linalg.cholesky).
144             The resulting lower-triangular matrix A is stored to facilitate the linear mapping W = A z
145             required for
146             path generation, as described in Section 2.1.
147
148             Parameters:
149             - S0: Initial stock price
150             - K: Strike price
151             - T: Time to maturity
152             - r: Risk-free interest rate
153             - sigma: Volatility
154             - d: Number of time steps (m in the explanation)
155             """
156             self.S0 = S0
157             self.K = K
158             self.T = T
159             self.r = r
160             self.sigma = sigma
161             self.d = d
162             self.dt = T / d # Time step size
163
164             # Covariance Matrix construction: _{ij} = min(t_i, t_j)
165             times = np.linspace(self.dt, T, d)
166             C = np.minimum(times[:, None], times[None, :])

```

```

165
166     # Cholesky Decomposition: A A^T =
167     self.A = np.linalg.cholesky(C)
168
169     # Precompute the drift term for each time step: (r - 0.5 ^2) * t_i
170     self.drift_term = (self.r - 0.5 * self.sigma**2) * times
171
172 def payoff(self, z):
173     """
174     payoff_arithmetic(z, model): Evaluates the discounted payoffs e^{-r T} _i(z).
175     Implements Eq. (4), returning e^{-r T} (S_mean - K)+.
176
177     Parameters:
178     - z: Array of standard normal random variables, shape (N, d)
179
180     Returns:
181     - Discounted payoffs, shape (N,)
182     """
183
184     if z.ndim == 1: z = z.reshape(1, -1) # Ensure 2D
185     # Correlated Brownian increments: W = A z
186     B = z @ self.A.T
187     # Stock prices: S_{t_i} = S0 * exp(drift + * W_{t_i})
188     S = self.S0 * np.exp(self.drift_term + self.sigma * B)
189     # Arithmetic average: (1/m) S_{t_i}
190     S_mean = np.mean(S, axis=1)
191     # Payoff: max(S_mean - K, 0), discounted
192     return np.exp(-self.r * self.T) * np.maximum(S_mean - self.K, 0)
193
194 # =====
195 # Directions & Shifts
196 # =====
197
198 def get_gradient(f, z, epsilon=1e-5):
199     d = len(z)
200     grad = np.zeros(d)
201     base = f(z)
202     for i in range(d):
203         z_p = z.copy()
204         z_p[i] += epsilon
205         grad[i] = (f(z_p) - base) / epsilon
206     return grad
207
208 def get_active_subspace(model, M=128):
209     """
210     get_active_subspace(model, payoff_fn, M): Approximates the solution to the maximization problem
211     (14).
212     It estimates the gradient covariance matrix C = E[      ^T] using a pilot sample of size M = 128.
213     Gradients are computed via finite differences (get_gradient). The function returns the leading
214     eigenvector
215     u1 of C, identifying the direction of the kink or jump.
216
217     Parameters:
218     - model: Option model instance
219     - M: Number of pilot samples (default 128)
220
221     Returns:
222     - u_as: Dominant direction vector (normalized)
223     """

```

```

221     sampler = qmc.Sobol(d=model.d, scramble=True) # Scrambled Sobol sequence
222     z_pilot = stats.norm.ppf(sampler.random(M)) # Transform to normal
223
224     grads = []
225     for i in range(M):
226         g = get_gradient(model.payoff, z_pilot[i]) # Finite difference gradient
227         grads.append(g)
228     grads = np.array(grads)
229
230     # Covariance matrix of gradients: C (1/M) ^T
231     C_hat = (grads.T @ grads) / M
232     evals, evecs = np.linalg.eigh(C_hat) # Eigen decomposition
233     u_as = evecs[:, -1] # Leading eigenvector (largest eigenvalue)
234     if np.sum(u_as) < 0: u_as = -u_as # Ensure consistent orientation
235     return u_as
236
237 def get_z1_direction(d):
238     """Direction for z1: [1, 0, ... 0]"""
239     u = np.zeros(d)
240     u[0] = 1.0
241     return u
242
243 def householder_matrix(u):
244     """
245     householder_matrix(u): Constructs the orthogonal matrix Q required to align the active direction
246     u1 with
247     the first canonical coordinate axis. It implements the Householder reflection Q = I - 2vv^T /
248     ||v||^2
249     with v = u - ||u|| e1, ensuring the numerically stable rotation of the integration domain.
250
251     Parameters:
252     - u: Input vector to rotate
253
254     Returns:
255     - Q: Orthogonal matrix
256     """
257     d = len(u)
258     e1 = np.zeros(d); e1[0] = 1.0 # First standard basis vector
259     sign = np.sign(u[0]) if u[0] != 0 else 1
260     v = u + sign * np.linalg.norm(u) * e1 # Householder vector
261     v = v / np.linalg.norm(v) # Normalize
262     H = np.eye(d) - 2 * np.outer(v, v) # Householder reflection
263     return H * (-sign) # Adjust sign
264
265 def get_odis_shift(model):
266     """
267     get_odis_shift_arithmetic(model): Solves the unconstrained optimization problem (23) for the
268     call option.
269     It minimizes the objective function J(z) = (1/2) ||z||^2 - ln((z)) using the L-BFGS-B algorithm.
270     To ensure global optimality for the non-convex objective, the optimizer is initialized from
271     multiple
272     starting points (including 0 and scaled vectors 1).
273
274     Parameters:
275     - model: Option model
276
277     Returns:
278     - : Optimal shift vector

```

```

275
276     """"
277     def obj(z):
278         p = model.payoff(z)[0]
279         if p <= 1e-12: return 1e6 # Penalty for zero payoff
280         return 0.5 * np.sum(z**2) - np.log(p) # Objective function J(z)
281
282     # Multi-start optimization to avoid local minima
283     best_res = None
284     best_val = np.inf
285     starts = [np.zeros(model.d), np.ones(model.d)*0.5, np.ones(model.d)*1.5] # Starting points
286
287     for x0 in starts:
288         res = minimize(obj, x0, method='L-BFGS-B') # Local optimization
289         if res.fun < best_val:
290             best_val = res.fun
291             best_res = res
292     return best_res.x # Optimal
293
294 # =====
295 # Estimators
296 # =====
297
298 def standard_estimator(model, N, method='MC', mu_shift=None):
299     """
300
301     standard_estimator(model, N, payoff_fn, method, ): A unified driver for the estimators V_{CMC}
302     and V_{QMC}.
303
304     Sequence Generation: If method='RQMC', it utilizes a scrambled Sobol sequence generator
305     (scipy.stats.qmc.Sobol)
306     to produce points U^{(i)} [0,1]^m, which are mapped to R^m via the inverse normal CDF ^{-1}.
307     To preserve numerical stability, inputs to ^{-1} are clipped to [10^{-10}, 1 - 10^{-10}].
308
309     Change of Measure: To support ODIS (Section 2.7), the function accepts an optimal drift
310     vector .
311
312     It shifts the samples Z' = Z + (simulating from q(z; )) and computes the Radon-Nikodym
313     derivative
314     w(z; ) = exp(- ^T Z' + (1/2)|| ||^2).
315
316     Estimation: Returns the sample mean of the product (Z') w(Z'; ), providing an unbiased
317     estimate of V.
318
319     Parameters:
320     - model: Option model
321     - N: Number of samples
322     - method: 'MC' for Monte Carlo, 'RQMC' for Quasi-Monte Carlo
323     - mu_shift: Shift vector for importance sampling (optional)
324
325     Returns:
326     - Estimated price (mean of weighted payoffs)
327     """
328
329     if method == 'MC':
330         z = np.random.standard_normal((N, model.d)) # Standard normal samples
331     elif method == 'RQMC':
332         sampler = qmc.Sobol(d=model.d, scramble=True) # Scrambled Sobol
333         u = sampler.random(N) # Uniform [0,1]^m
334         # Clip to avoid ^{-1} issues at boundaries
335         u = np.clip(u, 1e-10, 1 - 1e-10)
336         z = stats.norm.ppf(u) # Inverse CDF to normal
337
338     weights = np.ones(N) # Default weights
339     if mu_shift is not None:
340         z += mu_shift
341
342     payoff_fn = np.vectorize(lambda z: model.payoff(z)[0])
343     payoff_fn = np.broadcast_to(payoff_fn, (N, model.d))
344
345     payoff_fn = payoff_fn(z)
346
347     if method == 'MC':
348         payoff_fn *= weights
349         payoff_fn /= N
350
351     else:
352         payoff_fn *= weights
353         payoff_fn -= np.mean(weights)
354
355     return payoff_fn

```

```

328     # Importance Sampling: shift samples and adjust weights
329     X = z + mu_shift # Z' = Z +
330     dot = np.sum(X * mu_shift, axis=1) # ^T Z'
331     mu_sq = 0.5 * np.sum(mu_shift**2) # (1/2) |||^2
332     weights = np.exp(-dot + mu_sq) # w(Z'; )
333     payoffs = model.payoff(X) # (Z')
334 else:
335     payoffs = model.payoff(z) # (Z)
336
337 return np.mean(payoffs * weights) # Sample mean of w
338
339 def pre_int_estimator(model, N, u, Q, mu_perp=None, method='RQMC'):
340 """
341 pre_int_estimator(model, N, u, Q, _perp, method): Implements the pre-integrated estimator V_{PI}
342 for arithmetic Asian options (Section 2.8).
343     Pre-Integration: Smooths the payoff by analytically integrating over the last time step,
344     reducing variance for large m.
345     Closed-Form Smoothing: For arithmetic average A, the conditional expectation E[max(A - K, 0)
346     | Z_perp] is computed
347     using the closed-form expression (Eq. (25)), where A is approximated as a log-normal random
348     variable.
349     Active Subspace: Uses the active subspace direction u and rotation matrix Q to reduce
350     dimensionality.
351     Samples are generated in the perpendicular subspace Z_perp R^{m-1}.
352     Change of Measure: Applies importance sampling in the perpendicular subspace if _perp is
353     provided.
354     Estimation: For each sample, solves for v* such that the expected payoff equals K, then
355     computes the integral
356     using functions, providing a smoothed estimate of the option price.
357
358 Parameters:
359 - model: Option model
360 - N: Number of samples
361 - u: Active subspace direction vector
362 - Q: Rotation matrix from Householder transformation
363 - mu_perp: Shift vector in perpendicular subspace for importance sampling (optional)
364 - method: 'MC' or 'RQMC'
365
366 Returns:
367 - Estimated price (mean of smoothed payoffs)
368 """
369 d_perp = model.d - 1
370 if method == 'MC':
371     z_perp = np.random.standard_normal((N, d_perp))
372 else:
373     sampler = qmc.Sobol(d=d_perp, scramble=True)
374     z_perp = stats.norm.ppf(sampler.random(N))
375
376 weights = np.ones(N)
377 if mu_perp is not None:
378     X_perp = z_perp + mu_perp
379     dot = np.sum(X_perp * mu_perp, axis=1)
380     mu_sq = 0.5 * np.sum(mu_perp**2)
381     weights = np.exp(-dot + mu_sq)
382     z_perp = X_perp # Use shifted samples
383
384 # Geometric Constants
385 U_perp = Q[:, 1:]

```

```

379     Au = model.A @ u
380     AU_perp = model.A @ U_perp
381     beta = model.sigma * Au
382     const_part = np.log(model.S0) + model.drift_term
383
384     estimates = np.zeros(N)
385     exponent_perps = model.sigma * (z_perp @ AU_perp.T)
386
387     # We search for root in [-30, 30] to cover the relevant probability mass
388     for i in range(N):
389         alpha = np.exp(const_part + exponent_perps[i])
390         def g(v): return np.mean(alpha * np.exp(beta * v)) - model.K
391
392         # Root Finding
393         try: v_star = brentq(g, -30, 30)
394         except ValueError: v_star = -30 if g(0) > 0 else 30
395
396         # Closed-Form Integration
397         if v_star < 25:
398             d1 = beta - v_star
399             term1 = np.mean(alpha * np.exp(0.5 * beta**2) * stats.norm.cdf(d1))
400             term2 = model.K * stats.norm.cdf(-v_star)
401             val = np.exp(-model.r * model.T) * (term1 - term2)
402         else:
403             val = 0.0
404         estimates[i] = val * weights[i]
405
406     return np.mean(estimates)
407
408 # =====
409 # Simulation Logic
410 # =====
411
412 def run_experiment(K_target, d_val):
413     """
414     run_experiment(K_target, d_val): Runs the full simulation experiment for a given strike K and
415     dimension d.
416     Model Setup: Creates an AsianOptionCholesky model with the specified parameters.
417     Directions: Computes the z1 direction (first coordinate) and active subspace direction u_as,
418     along with their
419     rotation matrices Q_z1 and Q_as.
420     ODIS Shift: For K=120 (out-of-the-money), computes the optimal shift _opt using
421     get_odis_shift and projects
422     it to the perpendicular subspace for active subspace methods.
423     Ground Truth: Generates a high-accuracy reference value using pre_int_estimator with 2^17
424     samples.
425     Methods: Defines a dictionary of estimators to test, varying by K (K=100: basic methods;
426     K=120: includes ODIS).
427     Simulation Loop: For each sample size N and method, repeats N_REPEATS times to compute RMSE
428     and average time.
429     Output: Returns a DataFrame with results for plotting and analysis.
430
431     Parameters:
432     - K_target: Strike price (100 or 120)
433     - d_val: Dimension (number of time steps)
434
435     Returns:
436     - DataFrame with columns: K, N, Method, RMSE, Time, d

```

```

431     """
432     print(f"\n--- Running Experiment for K = {K_target}, d = {d_val} ---")
433     model = AsianOptionCholesky(S0=S0, K=K_target, T=T, r=r, sigma=sigma, d=d_val)
434
435     # Directions
436     u_z1 = get_z1_direction(d_val)
437     Q_z1 = np.eye(d_val)
438     u_as = get_active_subspace(model)
439     Q_as = householder_matrix(u_as)
440
441     # ODIS Shift (Only needed for K=120)
442     mu_opt = None
443     mu_perp_as = None
444     if K_target == 120:
445         mu_opt = get_odis_shift(model)
446         mu_local = Q_as.T @ mu_opt
447         mu_perp_as = mu_local[1:]
448
449     # Ground Truth Generation
450     if K_target == 120:
451         true_val = pre_int_estimator(model, 2**17, u_as, Q_as, mu_perp_as, 'RQMC')
452     else:
453         true_val = pre_int_estimator(model, 2**17, u_as, Q_as, None, 'RQMC')
454     print(f" -> Truth: {true_val:.6f}")
455
456     methods = {}
457     if K_target == 100:
458         methods['Crude MC'] = lambda n: standard_estimator(model, n, 'MC')
459         methods['Plain RQMC'] = lambda n: standard_estimator(model, n, 'RQMC')
460         methods['Pre-Int (z1)'] = lambda n: pre_int_estimator(model, n, u_z1, Q_z1, None, 'RQMC')
461         methods['Pre-Int (AS)'] = lambda n: pre_int_estimator(model, n, u_as, Q_as, None, 'RQMC')
462
463     elif K_target == 120:
464         methods['Crude MC'] = lambda n: standard_estimator(model, n, 'MC')
465         methods['MC + ODIS'] = lambda n: standard_estimator(model, n, 'MC', mu_opt)
466         methods['Plain RQMC'] = lambda n: standard_estimator(model, n, 'RQMC')
467         methods['RQMC + ODIS'] = lambda n: standard_estimator(model, n, 'RQMC', mu_opt)
468         methods['Pre-Int (AS)'] = lambda n: pre_int_estimator(model, n, u_as, Q_as, None, 'RQMC')
469         methods['Pre-Int (AS) + ODIS'] = lambda n: pre_int_estimator(model, n, u_as, Q_as,
470                                         mu_perp_as, 'RQMC')
471
472     results = []
473     total_ops = len(SAMPLE_SIZES) * len(methods) * N_REPEATS
474
475     with tqdm(total=total_ops, desc=f"Simulating d={d_val}") as pbar:
476         for N in SAMPLE_SIZES:
477             for name, func in methods.items():
478                 errs = []
479                 times = []
480                 for _ in range(N_REPEATS):
481                     t0 = time.time()
482                     est = func(N)
483                     t1 = time.time()
484                     errs.append(est)
485                     times.append(t1 - t0)
486                 pbar.update(1)
487
488             rmse = np.sqrt(np.mean((np.array(errs) - true_val)**2))

```

```

488         avg_time = np.mean(times)
489         results.append({'K': K_target, 'N': N, 'Method': name, 'RMSE': rmse, 'Time':
490             avg_time, 'd': d_val})
491
492     return pd.DataFrame(results)
493
494 # =====
495 # Plotting Functions
496 # =====
497 def get_convergence_rate(N, RMSE):
498     # Fit log(RMSE) = a + b * log(N)
499     # Slope b is the convergence rate
500     slope, intercept = np.polyfit(np.log(N), np.log(RMSE), 1)
501     return slope
502
503 def plot_k100(df, d_val, save_dir):
504     """
505     plot_k100(df, d_val, save_dir): Generates plots for K=100 experiments.
506         Convergence Plot (Left): Log-log plot of RMSE vs N for each method, with fitted convergence
507         rates.
508         Includes reference lines for O(N^{-0.5}) (MC) and O(N^{-1}) (QMC).
509         Efficiency Plot (Right): Log-log plot of RMSE vs computation time to assess
510         cost-effectiveness.
511         Output: Saves the figure as PNG in the specified directory.
512
513     Parameters:
514     - df: DataFrame with simulation results
515     - d_val: Dimension
516     - save_dir: Directory to save the plot
517     """
518
519     df = df[(df['K'] == 100) & (df['d'] == d_val)]
520     if df.empty: return
521
522     fig, axes = plt.subplots(1, 2, figsize=(16, 6))
523
524     # Convergence Plot
525     ax = axes[0]
526     methods = ['Crude MC', 'Plain RQMC', 'Pre-Int (z1)', 'Pre-Int (AS)']
527     markers = ['o', 's', 'x', '^']
528
529     for m, mark in zip(methods, markers):
530         sub = df[df['Method'] == m]
531         if sub.empty: continue
532
533         # Calculate Slope
534         slope = get_convergence_rate(sub['N'], sub['RMSE'])
535         label_str = f"{m} (Rate $\backslash\approx N^{{{{slope:.2f}}}}$)"
536
537         ax.loglog(sub['N'], sub['RMSE'], marker=mark, linestyle='-', label=label_str, base=2)
538
539     # Reference Lines
540     Ns = df['N'].unique()
541     ref_mc = Ns**(-0.5) * (df[df['Method']=='Crude MC']['RMSE'].iloc[0] * Ns[0]**0.5)
542     ax.loglog(Ns, ref_mc, 'k--', alpha=0.3, label='$O(N^{-0.5})$', base=2)
543
544     ref_qmc = Ns**(-1.0) * (df[df['Method']=='Plain RQMC']['RMSE'].iloc[0] * Ns[0]**1.0)
545     ax.loglog(Ns, ref_qmc, 'k:', alpha=0.3, label='$O(N^{-1.0})$', base=2)

```

```

543
544     ax.set_title(f'K=100, d={d_val}: Convergence Analysis', fontsize=14)
545     ax.set_xlabel('Sample Size $N$ (log2 scale)', fontsize=12)
546     ax.set_ylabel('RMSE', fontsize=12)
547     ax.legend(fontsize=10)
548     ax.grid(True, which="both", ls="--", alpha=0.2)
549     ax.xaxis.set_major_formatter(mticker.ScalarFormatter())
550
551     # 2. Computational Cost Plot
552     ax = axes[1]
553     for m, mark in zip(methods, markers):
554         sub = df[df['Method'] == m]
555         if sub.empty: continue
556         # Total time for N samples vs RMSE
557         ax.loglog(sub['Time'], sub['RMSE'], marker=mark, linestyle='--', label=m)
558
559     ax.set_title(f'K=100, d={d_val}: Efficiency (Error vs Time)', fontsize=14)
560     ax.set_xlabel('Avg Computation Time (s)', fontsize=12)
561     ax.set_ylabel('RMSE', fontsize=12)
562     ax.legend(fontsize=10)
563     ax.grid(True, which="both", ls="--", alpha=0.2)
564
565     plt.tight_layout()
566     fname = os.path.join(save_dir, f'Arithmetic_Asian_K100_d{d_val}_Analysis.png')
567     plt.savefig(fname, dpi=300)
568     print(f"Saved {fname}")
569     plt.close()
570
571 def plot_k120(df, d_val, save_dir):
572 """
573 plot_k120(df, d_val, save_dir): Generates plots for K=120 experiments (out-of-the-money case).
574     Plot A: Comprehensive comparison of all methods with convergence rates.
575     Plot B: Focus on variance reduction by comparing base methods vs methods with ODIS.
576     Pairs: Crude MC vs MC+ODIS, Plain RQMC vs RQMC+ODIS, Pre-Int (AS) vs Pre-Int (AS)+ODIS.
577     Output: Saves two PNG figures in the specified directory.
578
579 Parameters:
580 - df: DataFrame with simulation results
581 - d_val: Dimension
582 - save_dir: Directory to save the plots
583 """
584 df = df[(df['K'] == 120) & (df['d'] == d_val)]
585 if df.empty: return
586
587 # Plot A: Comprehensive Comparison
588 plt.figure(figsize=(10, 7))
589 methods = df['Method'].unique()
590
591 for m in methods:
592     sub = df[df['Method'] == m]
593     slope = get_convergence_rate(sub['N'], sub['RMSE'])
594     label_str = f"{{m}} (${{N}}^{{{slope:.2f}}})"
595     plt.loglog(sub['N'], sub['RMSE'], marker='o', label=label_str, base=2)
596
597 plt.title(f'K=120, d={d_val}: Comprehensive Comparison', fontsize=14)
598 plt.xlabel('Sample Size $N$ (log2)', fontsize=12)
599 plt.ylabel('RMSE', fontsize=12)
600 plt.grid(True, which="both", alpha=0.2)

```

```

601     plt.legend()
602     fname_A = os.path.join(save_dir, f'Arithmetic_Asian_K120_d{d_val}_Comprehensive.png')
603     plt.savefig(fname_A, dpi=300)
604     print(f"Saved {fname_A}")
605     plt.close()
606
607     # Plot B: Variance Reduction Focus
608     plt.figure(figsize=(10, 7))
609
610     pairs = [
611         ('Crude MC', 'MC + ODIS', 'red'),
612         ('Plain RQMC', 'RQMC + ODIS', 'blue'),
613         ('Pre-Int (AS)', 'Pre-Int (AS) + ODIS', 'green')
614     ]
615
616     for base, odis, color in pairs:
617         sub_b = df[df['Method'] == base]
618         if not sub_b.empty:
619             plt.loglog(sub_b['N'], sub_b['RMSE'], color=color, linestyle='--', marker='o',
620             label=base, base=2, alpha=0.5)
621
622         sub_o = df[df['Method'] == odis]
623         if not sub_o.empty:
624             slope = get_convergence_rate(sub_o['N'], sub_o['RMSE'])
625             label_str = f"{odis} ($N^{{{slope:.2f}}})"
626             plt.loglog(sub_o['N'], sub_o['RMSE'], color=color, linestyle='-', marker='D',
627             label=label_str, base=2)
628
629     plt.title(f'K=120, d={d_val}: Impact of Variance Reduction', fontsize=14)
630     plt.xlabel('Sample Size $N$ (log2)', fontsize=12)
631     plt.ylabel('RMSE', fontsize=12)
632     plt.grid(True, which="both", alpha=0.2)
633     plt.legend()
634     fname_B = os.path.join(save_dir, f'Arithmetic_Asian_K120_d{d_val}_Variance.png')
635     plt.savefig(fname_B, dpi=300)
636     print(f"Saved {fname_B}")
637     plt.close()
638
639 # =====
640 # Main Execution
641 # =====
642
643 if __name__ == "__main__":
644     """
645         Main execution block: Runs the full analysis for arithmetic Asian options.
646         Directory Setup: Creates 'plots_arithmetic_asian' and subdirectories for each dimension.
647         Experiment Loop: For each dimension d in DIMENSIONS, runs experiments for K=100 and K=120,
648         saves results to CSV, generates plots, and computes required N for K=120 with tolerance 0.01.
649         Output: Saves plots, CSVs, and required N files; prints summary of required N across
650         dimensions.
651     """
652
653     if not os.path.exists('plots_arithmetic_asian'):
654         os.makedirs('plots_arithmetic_asian')
655
656     required_n_results = []
657
658     # Loop over the dimensions

```

```
656     for d_val in DIMENSIONS:  
657         print("\n" + "#" * 60)  
658         print(f"PROCESSING DIMENSION: {d_val}")  
659         print("#" * 60)  
660  
661         # Create directory for this dimension  
662         curr_dir = os.path.join('plots_arithmetic_asian', f'd_{d_val}')  
663         if not os.path.exists(curr_dir):  
664             os.makedirs(curr_dir)  
665  
666         # Run Experiments  
667         df100 = run_experiment(100, d_val)  
668         df120 = run_experiment(120, d_val)  
669  
670         full_df = pd.concat([df100, df120])  
671         csv_name = os.path.join(curr_dir, f'arithmetic_asian_results_d{d_val}.csv')  
672         full_df.to_csv(csv_name, index=False)  
673         print(f"\nResults for d={d_val} saved to {csv_name}.")  
674  
675         # Generate Plots  
676         plot_k100(full_df, d_val, curr_dir)  
677         plot_k120(full_df, d_val, curr_dir)  
678  
679         # Find N for K=120  
680         print(f"\nFind N (Arithmetic) for d={d_val}")  
681         model_test = AsianOptionCholesky(S0=S0, K=120, T=T, r=r, sigma=sigma, d=d_val)  
682  
683         final_N, final_price = find_N_specific_arithmetic(model_test, d=d_val, tol=0.01)  
684         required_n_results.append({'d': d_val, 'Required_N': final_N, 'Estimated_Price':  
685             final_price})  
686  
687         with open(os.path.join(curr_dir, f'required_N_arithmetic_d{d_val}.txt'), 'w') as f:  
688             f.write(str(final_N))  
689  
690         print("\nArithmetic Option Analysis Complete..")
```

A.2 Digital Asian Option Implementation

```

1 import os
2 import time
3 import numpy as np
4 import pandas as pd
5 import scipy.stats as stats
6 from scipy.stats import qmc
7 from scipy.optimize import minimize, brentq
8 import matplotlib.pyplot as plt
9 import matplotlib.ticker as mticker
10 from tqdm import tqdm
11
12 np.random.seed(42)
13
14 """
15 This module implements Monte Carlo and Quasi-Monte Carlo simulations for pricing digital Asian
16 options.
17
18 Digital Asian options have binary payoffs based on whether the arithmetic average of the underlying
19 asset prices
20 exceeds a strike price. The code uses various variance reduction techniques including Quasi-Monte
21 Carlo (RQMC),
22 importance sampling (ODIS), pre-integration with closed-form smoothing, and active subspaces.
23
24 The simulations are performed for different dimensions (number of time steps) and strike prices,
25 with convergence analysis and efficiency comparisons.
26 """
27
28 # =====
29 # Setup and Configuration
30 # =====
31
32 # Model Parameters
33 S0 = 100 # Initial stock price
34 T = 1.0 # Time to maturity (in years)
35 r = 0.1 # Risk-free interest rate
36 sigma = 0.1 # Volatility of the stock
37
38 # Dimensions to test (number of time steps in the Asian option)
39 DIMENSIONS = [32, 64, 128, 256, 512]
40
41 # Simulation Parameters
42 POWERS = np.arange(7, 14) # Powers from 7 to 13, so sample sizes 2^7 to 2^13
43 SAMPLE_SIZES = 2**POWERS # Actual sample sizes: 128, 256, ..., 8192
44 N_REPEATS = 50 # Number of repetitions for each simulation to compute stable RMSE
45
46 def find_N_specific_digital(model, d, tol=0.01):
47 """
48 Implements the stopping criterion derived in Eq. (19). Rather than fixing N a priori, this
49 routine
50 increases the sample size in batches, monitoring the standard error of the estimator. It
51 terminates
52 the simulation only when the relative RMSE falls below the specified tolerance (e.g., =
53 10^{-2}),
54 ensuring the reliability of the OTM price estimates.
55
56 Parameters:

```

```

51     - model: An instance of DigitalAsianOption representing the option model.
52     - d: Dimension (number of time steps).
53     - tol: Tolerance for relative error (default 0.01, i.e., 1%).
54
55 Returns:
56 - N: The required sample size.
57 - mu: The estimated mean payoff.
58 """
59 z_score = 1.96 # 95% Confidence interval z-score
60
61 N = 128 # Starting sample size
62 # Generate initial batch of standard normal random variables
63 z_init = np.random.standard_normal((N, model.d))
64 Z_values = model.payoff_digital(z_init) # Compute payoffs for initial batch
65
66 # Initial statistics: mean and variance
67 mu = np.mean(Z_values)
68 var = np.var(Z_values, ddof=1) # Sample variance with ddof=1
69
70 print(f"\n--- Specific Recursive Search (Digital, d={d}, Start N={N}) ---")
71 print(f"{'N':<10} | {'Mean':<10} | {'Var':<10} | {'RelErr':<10}")
72 print("-" * 55)
73
74 while True:
75     # Generate one new standard normal sample
76     z_new = np.random.standard_normal((1, model.d))
77     Z_next = model.payoff_digital(z_new)[0] # Payoff for the new sample
78
79     mu_old = mu # Save previous mean
80
81     # Recursive update of mean and variance
82     mu = (N / (N + 1)) * mu + Z_next / (N + 1)
83     var = ((N - 1) / N) * var + (Z_next - mu_old)**2 / (N + 1)
84
85     N += 1 # Increment sample size
86
87     # Check for convergence
88     sigma_est = np.sqrt(max(var, 0.0)) # Estimated standard deviation
89
90     if abs(mu) > 1e-12 and sigma_est > 0:
91         half_width = z_score * sigma_est / np.sqrt(N) # Half-width of confidence interval
92         rel_err = half_width / abs(mu) # Relative error
93
94         if rel_err <= tol:
95             print("-" * 55)
96             print(f"Converged at N = {N} for d={d}")
97             print(f"Final Mean: {mu:.6f}")
98             print(f"Final Var: {var:.6f}")
99             return N, mu
100
101     # Print progress every 5000 samples
102     if N % 5000 == 0:
103         print(f"{'N':<10} | {'mu:<10.4f'} | {'var:<10.4f'} | {'rel_err:<10.2%'}")
104
105 # =====
106 # Model: Digital Asian Option
107 # =====
108

```

```

109 class DigitalAsianOption:
110     """
111     Represents a digital Asian option using Cholesky decomposition for path generation.
112
113     This class models the underlying asset paths using correlated Brownian motions,
114     discretized into 'd' time steps. The payoff is binary based on the arithmetic average
115     of the stock prices exceeding the strike.
116
117     Attributes:
118     - S0: Initial stock price
119     - K: Strike price
120     - T: Time to maturity
121     - r: Risk-free rate
122     - sigma: Volatility
123     - d: Number of time steps
124     - dt: Time step size
125     - A: Cholesky matrix for correlation
126     - drift_term: Precomputed drift term for each time step
127     """
128     def __init__(self, S0=100, K=100, T=1.0, r=0.1, sigma=0.1, d=32):
129         """
130         create_model(S0, K, T, r, , m): This initialization routine constructs the discrete time grid
131         t_i = i*T/m where i = 1, ..., m and the associated covariance matrix R^{mm} with entries
132         _{ij} = min(t_i, t_j). Crucially, it performs the Cholesky decomposition = A A^T (via
133         numpy.linalg.cholesky).
134         The resulting lower-triangular matrix A is stored to facilitate the linear mapping W = A z
135         required for
136             path generation, as described in Section 2.1.
137
138         Parameters:
139             - S0: Initial stock price
140             - K: Strike price
141             - T: Time to maturity
142             - r: Risk-free interest rate
143             - sigma: Volatility
144             - d: Number of time steps (m in the explanation)
145
146         self.S0 = S0
147         self.K = K
148         self.T = T
149         self.r = r
150         self.sigma = sigma
151         self.d = d
152         self.dt = T / d # Time step size
153
154         # 1. Covariance Matrix construction: _{ij} = min(t_i, t_j)
155         times = np.linspace(self.dt, T, d)
156         C = np.minimum(times[:, None], times[None, :])
157
158         # Cholesky Decomposition: A A^T =
159         self.A = np.linalg.cholesky(C)
160
161         # Precompute drift for simulation: (r - 0.5 ^2) * t_i
162         self.drift_term = (self.r - 0.5 * self.sigma**2) * times
163
164     def get_S_mean(self, z):
165         """
166         get_s_mean(z, model): Computes the realization of the random variable (1/m) _{i=1}^m S_{t_i}.

```

```

165     Given a standard normal vector  $z \in \mathbb{R}^m$ , it first recovers the Brownian motion path  $W = A z$ .
166     It then applies the geometric Brownian motion mapping  $S_{\{t_i\}} = S_0 \exp((r - \sigma^2/2) t_i +$ 
167      $W_{\{t_i\}})$ 
168     in a vectorized manner. This function serves as the numerical evaluation of the underlying
169     asset process.
170
171     Parameters:
172     -  $z$ : Array of standard normal random variables, shape (N, d)
173
174     Returns:
175     - Arithmetic mean of stock prices, shape (N,)
176     """
177     if z.ndim == 1: z = z.reshape(1, -1) # Ensure 2D
178     # Correlated Brownian increments:  $W = A z$ 
179     B = z @ self.A.T
180     # Stock prices:  $S_{\{t_i\}} = S_0 * \exp(\text{drift} + \sigma * W_{\{t_i\}})$ 
181     S = self.S0 * np.exp(self.drift_term + self.sigma * B)
182     # Arithmetic mean:  $(1/m) \sum S_{\{t_i\}}$ 
183     return np.mean(S, axis=1)
184
185     def payoff_digital(self, z):
186         """
187         payoff_digital(z, model): Evaluates the discounted payoffs  $e^{-r T} \max(S_{\{t_i\}} - K, 0)$ .
188         Implements Eq. (5), returning  $e^{-r T}$  if the arithmetic mean exceeds K, and 0 otherwise.
189
190         Parameters:
191         -  $z$ : Array of standard normal random variables, shape (N, d)
192
193         Returns:
194         - Discounted binary payoffs, shape (N,)
195         """
196         S_mean = self.get_S_mean(z) # Compute arithmetic mean
197         p = np.where(S_mean > self.K, 1.0, 0.0) # Binary payoff
198         return np.exp(-self.r * self.T) * p # Discounted
199
200     def payoff_arithmetic(self, z):
201         """
202         payoff_arithmetic(z, model): Evaluates the discounted payoffs  $e^{-r T} \max(S_{\{t_i\}} - K, 0)$  for the
203         arithmetic option.
204         Returns  $e^{-r T} \max(S_{\{t_i\}} - K, 0)$ . Used as a proxy for gradient estimation in active
205         subspaces.
206
207         Parameters:
208         -  $z$ : Array of standard normal random variables, shape (N, d)
209
210         Returns:
211         - Discounted arithmetic payoffs, shape (N,)
212         """
213         S_mean = self.get_S_mean(z) # Compute arithmetic mean
214         return np.exp(-self.r * self.T) * np.maximum(S_mean - self.K, 0) # Discounted payoff
215
216     # =====
217     # Directions & Shifts
218     # =====
219
220     def get_gradient(f, z, epsilon=1e-5):
221         d = len(z)
222         grad = np.zeros(d)

```

```

219     base = f(z)
220     for i in range(d):
221         z_p = z.copy()
222         z_p[i] += epsilon
223         grad[i] = (f(z_p) - base) / epsilon
224     return grad
225
226 def get_active_subspace(model, M=128):
227     """
228     get_active_subspace(model, payoff_fn, M): Approximates the solution to the maximization problem
229     (14).
230     It estimates the gradient covariance matrix  $C = E[\nabla f]^T$  using a pilot sample of size  $M = 128$ .
231     Gradients are computed via finite differences (get_gradient). The function returns the leading
232     eigenvector
233     u1 of C, identifying the direction of the kink or jump.
234
235     Parameters:
236     - model: Option model instance
237     - M: Number of pilot samples (default 128)
238
239     Returns:
240     - u_as: Dominant direction vector (normalized)
241     """
242
243     sampler = qmc.Sobol(d=model.d, scramble=True) # Scrambled Sobol sequence
244     z_pilot = stats.norm.ppf(sampler.random(M)) # Transform to normal
245
246     grads = []
247     # Use Arithmetic payoff for gradient estimation (as proxy for digital)
248     for i in range(M):
249         g = get_gradient(model.payoff_arithmetic, z_pilot[i]) # Finite difference gradient
250         grads.append(g)
251     grads = np.array(grads)
252
253     # Covariance matrix of gradients:  $C = \frac{1}{M} \nabla f^T$ 
254     C_hat = (grads.T @ grads) / M
255     evals, evecs = np.linalg.eigh(C_hat) # Eigen decomposition
256     u_as = evecs[:, -1] # Leading eigenvector (largest eigenvalue)
257
258     # Standardize sign: ensure mean moves up with positive u
259     test_z = u_as * 0.1
260     if model.get_S_mean(test_z)[0] < model.get_S_mean(-test_z)[0]:
261         u_as = -u_as
262     return u_as
263
264 def get_z1_direction(d):
265     """Direction for z1: [1, 0, ... 0]"""
266     u = np.zeros(d)
267     u[0] = 1.0
268     return u
269
270 def householder_matrix(u):
271     """Orthogonal matrix Q where first column is u"""
272     d = len(u)
273     e1 = np.zeros(d); e1[0] = 1.0
274     sign = np.sign(u[0]) if u[0] != 0 else 1
275     v = u + sign * np.linalg.norm(u) * e1
276     v = v / np.linalg.norm(v)
277     H = np.eye(d) - 2 * np.outer(v, v)

```

```

275     return H * (-sign)
276
277 def get_odis_shift(model):
278     """
279     get_odis_shift_digital(model): Determines the optimal drift * for the binary option.
280     Instead of the generic likelihood maximization, it solves a constrained geometric problem:
281     it minimizes the distance ||z||^2 / 2 subject to (z) > 0. This identifies the point on the limit
282     surface (z) = 0 with the highest probability density, solved via the SLSQP algorithm.
283
284     Parameters:
285     - model: Option model
286
287     Returns:
288     - Optimal shift vector *
289     """
290
291     def constr(z):
292         # Constraint: S_mean(z) - K = 0 (on the exercise boundary)
293         return model.get_S_mean(z)[0] - model.K
294
295     def obj(z):
296         # Objective: minimize ||z||^2 / 2
297         return 0.5 * np.sum(z**2)
298
299     x0 = np.ones(model.d) * 0.1 # Initial guess
300     cons = ({'type': 'eq', 'fun': constr}) # Equality constraint
301     res = minimize(obj, x0, method='SLSQP', constraints=cons, tol=1e-4) # Constrained optimization
302
303     if not res.success:
304         return np.zeros(model.d) # Fallback to zero shift
305
306     return res.x # Optimal
307
308 # =====
309 # Estimators
310 # =====
311
312 def standard_estimator(model, N, method='MC', mu_shift=None):
313     """
314     standard_estimator(model, N, payoff_fn, method, ): A unified driver for the estimators V_{CMC}
315     and V_{QMC}.
316     Sequence Generation: If method='RQMC', it utilizes a scrambled Sobol sequence generator
317     (scipy.stats.qmc.Sobol)
318     to produce points U^{(i)} [0,1]^m, which are mapped to R^m via the inverse normal CDF ^{-1}.
319     To preserve numerical stability, inputs to ^{-1} are clipped to [10^{-10}, 1 - 10^{-10}].
320     Change of Measure: To support ODIS (Section 2.7), the function accepts an optimal drift
321     vector .
322     It shifts the samples Z' = Z + (simulating from q(z; )) and computes the Radon-Nikodym
323     derivative
324     w(z; ) = exp(- ^T Z' + (1/2)|||^2).
325     Estimation: Returns the sample mean of the product (Z') w(Z'; ), providing an unbiased
326     estimate of V.
327
328     Parameters:
329     - model: Option model
330     - N: Number of samples
331     - method: 'MC' for Monte Carlo, 'RQMC' for Quasi-Monte Carlo
332     - mu_shift: Shift vector for importance sampling (optional)

```

```

328     Returns:
329     - Estimated price (mean of weighted payoffs)
330     """
331     if method == 'MC':
332         z = np.random.standard_normal((N, model.d)) # Standard normal samples
333     elif method == 'RQMC':
334         sampler = qmc.Sobol(d=model.d, scramble=True) # Scrambled Sobol
335         u = sampler.random(N) # Uniform [0,1]^m
336         # Clip to avoid ^{-1} issues at boundaries
337         u = np.clip(u, 1e-10, 1 - 1e-10)
338         z = stats.norm.ppf(u) # Inverse CDF to normal
339
340     weights = np.ones(N) # Default weights
341     if mu_shift is not None:
342         # Importance Sampling: shift samples and adjust weights
343         X = z + mu_shift # Z' = Z +
344         dot = np.sum(X * mu_shift, axis=1) # ^T Z'
345         mu_sq = 0.5 * np.sum(mu_shift**2) # (1/2) ||| |^2
346         weights = np.exp(-dot + mu_sq) # w(Z'; )
347         payoffs = model.payoff_digital(X) # (Z')
348     else:
349         payoffs = model.payoff_digital(z) # (Z)
350
351     return np.mean(payoffs * weights) # Sample mean of w
352
353 def pre_int_estimator_closed_form(model, N, u, Q, mu_perp=None, method='RQMC'):
354     """
355     vector_pre_int_digital(...): These functions implement the smoothed estimators by evaluating the
356     inner
357     integral p(x_{-1}) defined in Eq. (6).
358     Root Finding: For every sample z_{-1} in the subspace R^{m-1}, the routine defines the
359     implicit function
360     g(v) = Mean(S(v u1 + Q_{-1} z_{-1})) - K. It employs Brent's method to find the critical value
361     v* = (z_{-1})
362     where the payoff activates.
363     Analytic Smoothing: For the Digital option, the discontinuous indicator is replaced by the
364     smooth tail
365     probability (-v*). For the Arithmetic option, the function computes the conditional expectation
366     E[(S_mean - K)^+ | z_{-1}] analytically using Gaussian integrals, removing the derivative
367     discontinuity.
368
369     Parameters:
370     - model: Option model
371     - N: Number of samples in perpendicular subspace
372     - u: Direction vector (e.g., active subspace or z1)
373     - Q: Orthogonal matrix from Householder
374     - mu_perp: Shift in perpendicular direction (optional)
375     - method: 'MC' or 'RQMC'
376
377     Returns:
378     - Estimated price using pre-integration
379     """
380     d_perp = model.d - 1 # Dimension of perpendicular subspace R^{m-1}
381     if method == 'MC':
382         z_perp = np.random.standard_normal((N, d_perp)) # Samples in R^{m-1}
383     else:
384         sampler = qmc.Sobol(d=d_perp, scramble=True)
385         z_perp = stats.norm.ppf(sampler.random(N)) # Quasi-random

```

```

381
382     weights = np.ones(N) # IS weights
383     if mu_perp is not None:
384         # Apply IS in perpendicular direction
385         X_perp = z_perp + mu_perp
386         dot = np.sum(X_perp * mu_perp, axis=1)
387         mu_sq = 0.5 * np.sum(mu_perp**2)
388         weights = np.exp(-dot + mu_sq)
389         z_perp = X_perp # Use shifted samples
390
391     # Precompute matrices
392     U_perp = Q[:, 1:] # Perpendicular directions
393     Au = model.A @ u # A u
394     AU_perp_z_perp = (model.A @ U_perp) @ z_perp.T # A U_{-1} z_{-1}
395
396     beta = model.sigma * Au # Coefficient for exponential
397     const_log_S = np.log(model.S0) + model.drift_term # Log S0 + drift terms
398
399     estimates = np.zeros(N)
400
401     # Search range for v*
402     BOUND = 30.0
403
404     for i in range(N):
405         # Contribution from perpendicular part
406         perp_contribution = model.sigma * AU_perp_z_perp[:, i]
407         alpha = np.exp(const_log_S + perp_contribution) # Alpha coefficients
408
409         def g(v):
410             # g(v) = E[S_mean | z_{-1}] - K, where S_mean depends on v
411             return np.mean(alpha * np.exp(beta * v)) - model.K
412
413         # Find v* such that g(v*) = 0
414         try:
415             if g(-BOUND) * g(BOUND) < 0:
416                 v_star = brentq(g, -BOUND, BOUND) # Root in [-BOUND, BOUND]
417             else:
418                 v_star = -BOUND if g(0) > 0 else BOUND # Boundary case
419         except ValueError:
420             v_star = -BOUND if g(0) > 0 else BOUND
421
422         # Closed-form smoothing: replace indicator with (-v*)
423         val = np.exp(-model.r * model.T) * stats.norm.cdf(-v_star)
424         estimates[i] = val * weights[i] # Weighted estimate
425
426     return np.mean(estimates) # Average over samples
427
428     # =====
429     # Simulation Logic
430     # =====
431
432     def run_experiment(K_target, d_val):
433         """
434             Runs simulations for a given strike price K_target and dimension d_val.
435
436             Computes RMSE for various estimators and returns results as a DataFrame.
437         """
438         print(f"\n--- Running Experiment for K = {K_target}, d = {d_val} (Digital) ---")

```

```

439     model = DigitalAsianOption(S0=S0, K=K_target, T=T, r=r, sigma=sigma, d=d_val) # Initialize model
440
441     # Compute directions for pre-integration
442     u_z1 = get_z1_direction(d_val) # z1 direction
443     Q_z1 = np.eye(d_val) # Identity for z1
444     u_as = get_active_subspace(model) # Active subspace direction
445     Q_as = householder_matrix(u_as) # Rotation matrix
446
447     # Compute ODIS shift only for OTM case (K=120)
448     mu_opt = None
449     mu_perp_as = None
450     if K_target == 120:
451         print(" -> Computing ODIS shift...")
452         mu_opt = get_odis_shift(model) # Optimal shift
453         mu_local = Q_as.T @ mu_opt # Transform to local coords
454         mu_perp_as = mu_local[1:] # Perpendicular component
455
456     # Compute ground truth using high-accuracy pre-integration
457     print(" -> Computing Ground Truth...")
458     if K_target == 120:
459         true_val = pre_int_estimator_closed_form(model, 2**17, u_as, Q_as, mu_perp_as, 'RQMC')
460     else:
461         true_val = pre_int_estimator_closed_form(model, 2**17, u_as, Q_as, None, 'RQMC')
462     print(f" -> Truth: {true_val:.6f}")
463
464     # Define methods based on strike
465     methods = {}
466     if K_target == 100: # ATM
467         methods['Crude MC'] = lambda n: standard_estimator(model, n, 'MC')
468         methods['Plain RQMC'] = lambda n: standard_estimator(model, n, 'RQMC')
469         methods['Pre-Int (z1)'] = lambda n: pre_int_estimator_closed_form(model, n, u_z1, Q_z1,
470             None, 'RQMC')
471         methods['Pre-Int (AS)'] = lambda n: pre_int_estimator_closed_form(model, n, u_as, Q_as,
472             None, 'RQMC')
473
474     elif K_target == 120: # OTM
475         methods['Crude MC'] = lambda n: standard_estimator(model, n, 'MC')
476         methods['MC + ODIS'] = lambda n: standard_estimator(model, n, 'MC', mu_opt)
477         methods['Plain RQMC'] = lambda n: standard_estimator(model, n, 'RQMC')
478         methods['RQMC + ODIS'] = lambda n: standard_estimator(model, n, 'RQMC', mu_opt)
479         methods['Pre-Int (AS)'] = lambda n: pre_int_estimator_closed_form(model, n, u_as, Q_as,
480             None, 'RQMC')
481         methods['Pre-Int (AS) + ODIS'] = lambda n: pre_int_estimator_closed_form(model, n, u_as,
482             Q_as, mu_perp_as, 'RQMC')
483
484     results = []
485     total_ops = len(SAMPLE_SIZES) * len(methods) * N_REPEATS # Progress tracking
486
487     with tqdm(total=total_ops, desc=f"Simulating d={d_val}") as pbar:
488         for N in SAMPLE_SIZES:
489             for name, func in methods.items():
490                 errs = []
491                 times = []
492                 for _ in range(N_REPEATS):
493                     t0 = time.time()
494                     est = func(N)
495                     t1 = time.time()
496                     errs.append(est)

```

```

493         times.append(t1 - t0)
494         pbar.update(1)
495
496         rmse = np.sqrt(np.mean((np.array(errs) - true_val)**2))
497         avg_time = np.mean(times)
498         results.append({'K': K_target, 'N': N, 'Method': name, 'RMSE': rmse, 'Time': avg_time, 'd': d_val})
499
500     return pd.DataFrame(results)
501
502 # =====
503 # Plotting Functions
504 # =====
505
506 def get_convergence_rate(N, RMSE):
507     slope, intercept = np.polyfit(np.log(N), np.log(RMSE), 1)
508     return slope
509
510 def plot_k100(df, d_val, save_dir):
511     """
512     K=100 Plot: Convergence and efficiency analysis for digital option.
513     Left: Convergence (Log2 N vs RMSE)
514     Right: Cost (Time vs RMSE)
515     """
516     df = df[(df['K'] == 100) & (df['d'] == d_val)]
517     if df.empty: return
518
519     fig, axes = plt.subplots(1, 2, figsize=(16, 6))
520
521     # Convergence Plot
522     ax = axes[0]
523     methods = ['Crude MC', 'Plain RQMC', 'Pre-Int (z1)', 'Pre-Int (AS)']
524     markers = ['o', 's', 'x', '^']
525
526     for m, mark in zip(methods, markers):
527         sub = df[df['Method'] == m]
528         if sub.empty: continue
529
530         slope = get_convergence_rate(sub['N'], sub['RMSE'])
531         label_str = f'{m} (Rate $\backslash\approx N^{{{slope:.2f}}})$'
532         ax.loglog(sub['N'], sub['RMSE'], marker=mark, linestyle='-', label=label_str, base=2)
533
534     Ns = df['N'].unique()
535     # Reference Lines
536     if not df[df['Method']=='Crude MC'].empty:
537         ref_mc = Ns**(-0.5) * (df[df['Method']=='Crude MC']['RMSE'].iloc[0] * Ns[0]**0.5)
538         ax.loglog(Ns, ref_mc, 'k--', alpha=0.3, label='$O(N^{-0.5})$', base=2)
539
540     if not df[df['Method']=='Plain RQMC'].empty:
541         ref_qmc = Ns**(-1.0) * (df[df['Method']=='Plain RQMC']['RMSE'].iloc[0] * Ns[0]**1.0)
542         ax.loglog(Ns, ref_qmc, 'k:', alpha=0.3, label='$O(N^{-1.0})$', base=2)
543
544     ax.set_title(f'K=100 (Digital No-Quad, d={d_val}): Convergence Analysis', fontsize=14)
545     ax.set_xlabel('Sample Size $N$ (log2 scale)', fontsize=12)
546     ax.set_ylabel('RMSE', fontsize=12)
547     ax.legend(fontsize=10)
548     ax.grid(True, which="both", ls="--", alpha=0.2)
549     ax.xaxis.set_major_formatter(mticker.ScalarFormatter())

```

```

550
551 # Computational Cost Plot
552 ax = axes[1]
553 for m, mark in zip(methods, markers):
554     sub = df[df['Method'] == m]
555     if sub.empty: continue
556     # Total time for N samples vs RMSE
557     ax.loglog(sub['Time'], sub['RMSE'], marker=mark, linestyle='-', label=m)
558
559     ax.set_title(f'K=100 (Digital No-Quad, d={d_val}): Efficiency', fontsize=14)
560     ax.set_xlabel('Avg Computation Time (s)', fontsize=12)
561     ax.set_ylabel('RMSE', fontsize=12)
562     ax.legend(fontsize=10)
563     ax.grid(True, which="both", ls="--", alpha=0.2)
564
565 plt.tight_layout()
566 fname = os.path.join(save_dir, f'Digital_NoQuad_K100_d{d_val}_Analysis.png')
567 plt.savefig(fname, dpi=300)
568 print(f"Saved {fname}")
569 plt.close()

570
571 def plot_k120(df, d_val, save_dir):
572 """
573 K=120 Plots: Comprehensive comparison and variance reduction focus for digital option.
574 Plot A: Comprehensive (All methods).
575 Plot B: Variance Reduction Focus (Method vs Method+ODIS).
576 """
577 df = df[(df['K'] == 120) & (df['d'] == d_val)]
578 if df.empty: return
579
580 # Plot A: Comprehensive Comparison
581 plt.figure(figsize=(10, 7))
582 methods = df['Method'].unique()
583
584 for m in methods:
585     sub = df[df['Method'] == m]
586     if sub.empty: continue
587     slope = get_convergence_rate(sub['N'], sub['RMSE'])
588     label_str = f'{m} ($N^{{{{slope:.2f}}}})$'
589     plt.loglog(sub['N'], sub['RMSE'], marker='o', label=label_str, base=2)
590
591 plt.title(f'K=120 (Digital No-Quad, d={d_val}): Comprehensive Comparison', fontsize=14)
592 plt.xlabel('Sample Size $N$ (log2)', fontsize=12)
593 plt.ylabel('RMSE', fontsize=12)
594 plt.grid(True, which="both", alpha=0.2)
595 plt.legend()
596 fname_A = os.path.join(save_dir, f'Digital_NoQuad_K120_d{d_val}_Comprehensive.png')
597 plt.savefig(fname_A, dpi=300)
598 print(f"Saved {fname_A}")
599 plt.close()

600
601 # Plot B: Variance Reduction Focus
602 plt.figure(figsize=(10, 7))

603 pairs = [
604     ('Crude MC', 'MC + ODIS', 'red'),
605     ('Plain RQMC', 'RQMC + ODIS', 'blue'),
606     ('Pre-Int (AS)', 'Pre-Int (AS) + ODIS', 'green')

```

```

608     ]
609
610     for base, odis, color in pairs:
611         sub_b = df[df['Method'] == base]
612         if not sub_b.empty:
613             plt.loglog(sub_b['N'], sub_b['RMSE'], color=color, linestyle='--', marker='o',
614             label=base, base=2, alpha=0.5)
615
616         sub_o = df[df['Method'] == odis]
617         if not sub_o.empty:
618             slope = get_convergence_rate(sub_o['N'], sub_o['RMSE'])
619             label_str = f"{odis} ($N^{{{slope:.2f}}})$"
620             plt.loglog(sub_o['N'], sub_o['RMSE'], color=color, linestyle='-', marker='D',
621             label=label_str, base=2)
622
623     plt.title(f'K=120 (Digital No-Quad, d={d_val}): Variance Reduction Impact', fontsize=14)
624     plt.xlabel('Sample Size $N$ (log2)', fontsize=12)
625     plt.ylabel('RMSE', fontsize=12)
626     plt.grid(True, which="both", alpha=0.2)
627     plt.legend()
628     fname_B = os.path.join(save_dir, f'Digital_NoQuad_K120_d{d_val}_Variance.png')
629     plt.savefig(fname_B, dpi=300)
630     print(f"Saved {fname_B}")
631     plt.close()
632
633 # =====
634 # Main Execution
635 # =====
636
637 if __name__ == "__main__":
638     # Main execution for digital Asian option analysis
639     if not os.path.exists('plots_digital_asian_2'):
640         os.makedirs('plots_digital_asian_2')
641
642     required_n_results = []
643
644     # Loop over the dimensions
645     for d_val in DIMENSIONS:
646         print("\n" + "#" * 60)
647         print(f"PROCESSING DIMENSION: {d_val}")
648         print("#" * 60)
649
650         # Create directory for this dimension
651         curr_dir = os.path.join('plots_digital_asian_2', f'd_{d_val}')
652         if not os.path.exists(curr_dir):
653             os.makedirs(curr_dir)
654
655         # Run Experiments
656         df100 = run_experiment(100, d_val)
657         df120 = run_experiment(120, d_val)
658
659         full_df = pd.concat([df100, df120])
660         csv_name = os.path.join(curr_dir, f'digital_option_results_2_d{d_val}.csv')
661         full_df.to_csv(csv_name, index=False)
662         print(f"\nResults for d={d_val} saved to {csv_name}.")
663
664         # Generate Plots
665         plot_k100(full_df, d_val, curr_dir)

```

```
664     plot_k120(full_df, d_val, curr_dir)
665
666     # Find N for K=120
667     print(f"\nFind N (Digital) for d={d_val}")
668     model_test = DigitalAsianOption(S0=S0, K=120, T=T, r=r, sigma=sigma, d=d_val)
669
670     final_N, final_price = find_N_specific_digital(model_test, d=d_val, tol=0.01)
671     required_n_results.append({'d': d_val, 'Required_N': final_N, 'Estimated_Price': final_price})
672
673     with open(os.path.join(curr_dir, f'required_N_digital_d{d_val}.txt'), 'w') as f:
674         f.write(str(final_N))
675
676 print("\nDigital Option Analysis Complete.")
```

AI Use Disclosure In compliance with point 3 of the project rules, we disclose the usage of AI assistance in the development of this project. AI tools were utilized for the following specific tasks:

- **Code Optimization and Plotting:** The AI assisted in organizing the Python code structure for better modularity and readability, as well as generating the scripts used to produce the convergence and efficiency plots presented in the report.
- **Active Subspace Implementation for Digital Options:** The AI provided guidance on handling the lack of useful gradients in the Digital Asian option payoff (which is piecewise constant). Specifically, it helped implement the strategy of using the Arithmetic Asian option as a smooth proxy to estimate the gradient covariance matrix and identify the active subspace directions effectively.